

CS 106B

Lecture 27: A* Heuristics and Minimum Spanning Trees

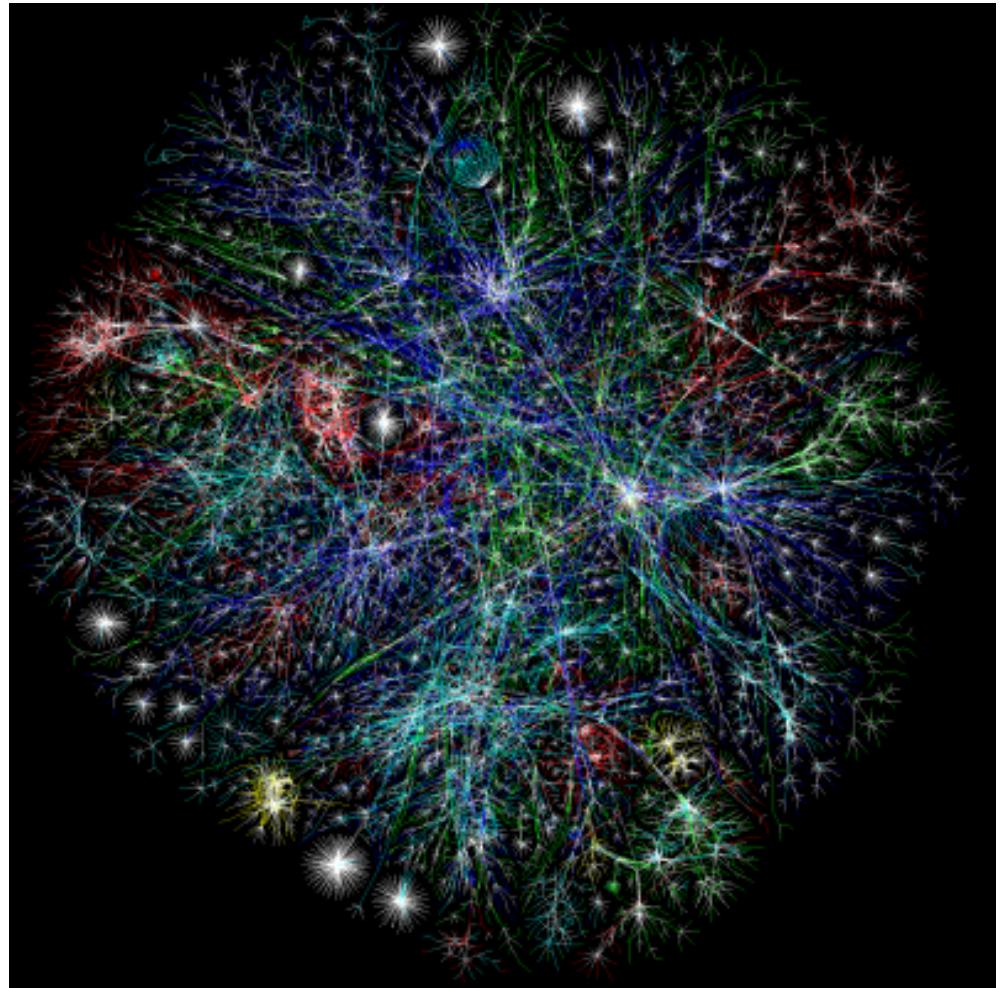
Friday, December 2, 2016

Programming Abstractions
Fall 2016
Stanford University
Computer Science Department

Lecturer: Chris Gregg

reading:

Programming Abstractions in C++, pp. 820-821



Today's Topics

- Logistics
 - We updated the Trailblazer alternate route description on the handout.
- Typedef definition
- Real Graph: Internet routers and traceroute
- More on Trailblazer
- A* Heuristics
 - heuristic bounds, and why we want to underestimate
- Minimum Spanning Trees
 - Kruskal's algorithm
- Dijkstra and negative weights (extra slides)



C++: `typedef`

In Assignment 7, you will use a `Vector<Vertex *>` type for all your paths. Because paths are used so often in the assignment, we have defined a type called `Path`, which is simply a `Vector<Vertex *>`. To define a new type in C++, you can use the `typedef` keyword, as follows:

```
typedef Vector<Vertex *> Path;
```

Now, you can use `Path` as any other variable, and it is just a `Vector<Vertex *>`:

```
Path p;
if (p.size() > 0) {
    Vertex *v = p[0];
}
for (Vertex *v : p) {
    cout << "Next vertex name: " << v->name << endl;
}
```



Real Graphs!

There was a Tiny Feedback from the last lecture that said,

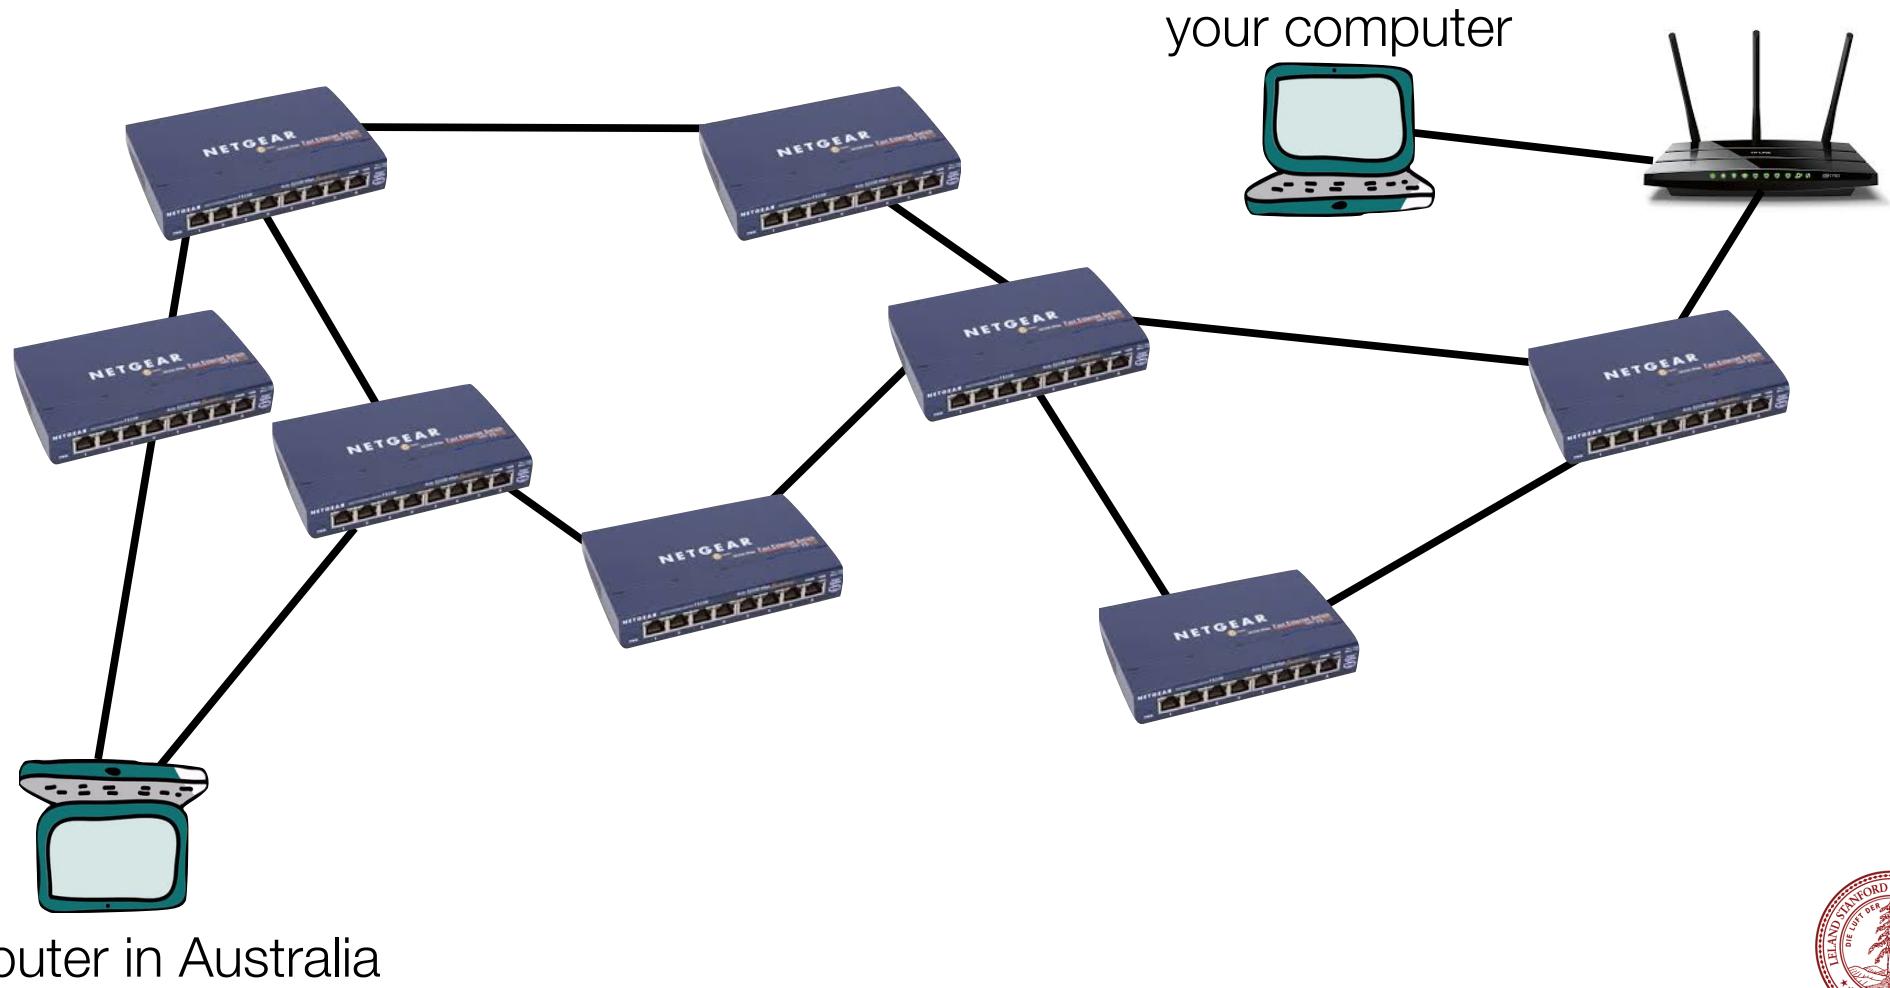
“Would love more stories about when you two use these different search algorithms more in real life.”

On Monday we played the "Wikipedia Get to Philosophy" game with the Internet, which have web pages with links that form a graph. Let's see another example of how the Internet is a real graph in a completely different way: Routers

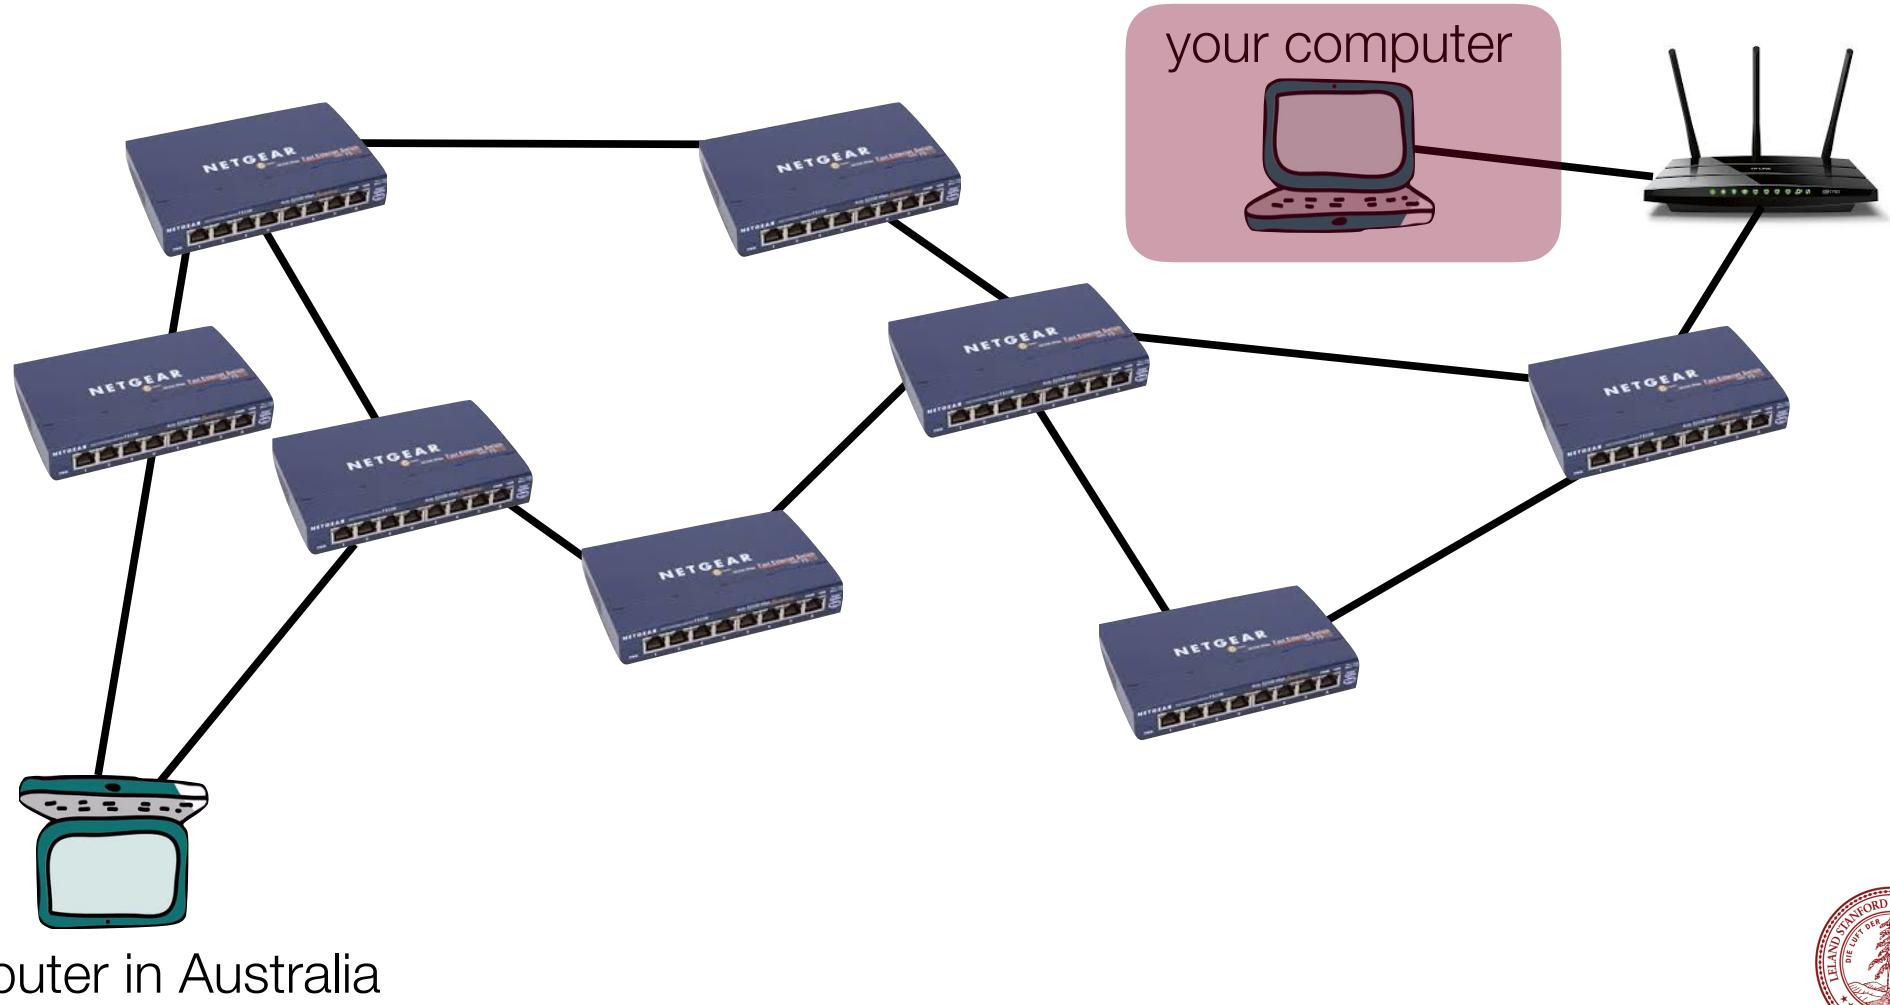
How does a message get sent from your computer to another computer on the Internet, say in Australia?



The Internet: Computers connected through routers



The Internet: Computers connected through routers



The Internet: Let's simplify a bit

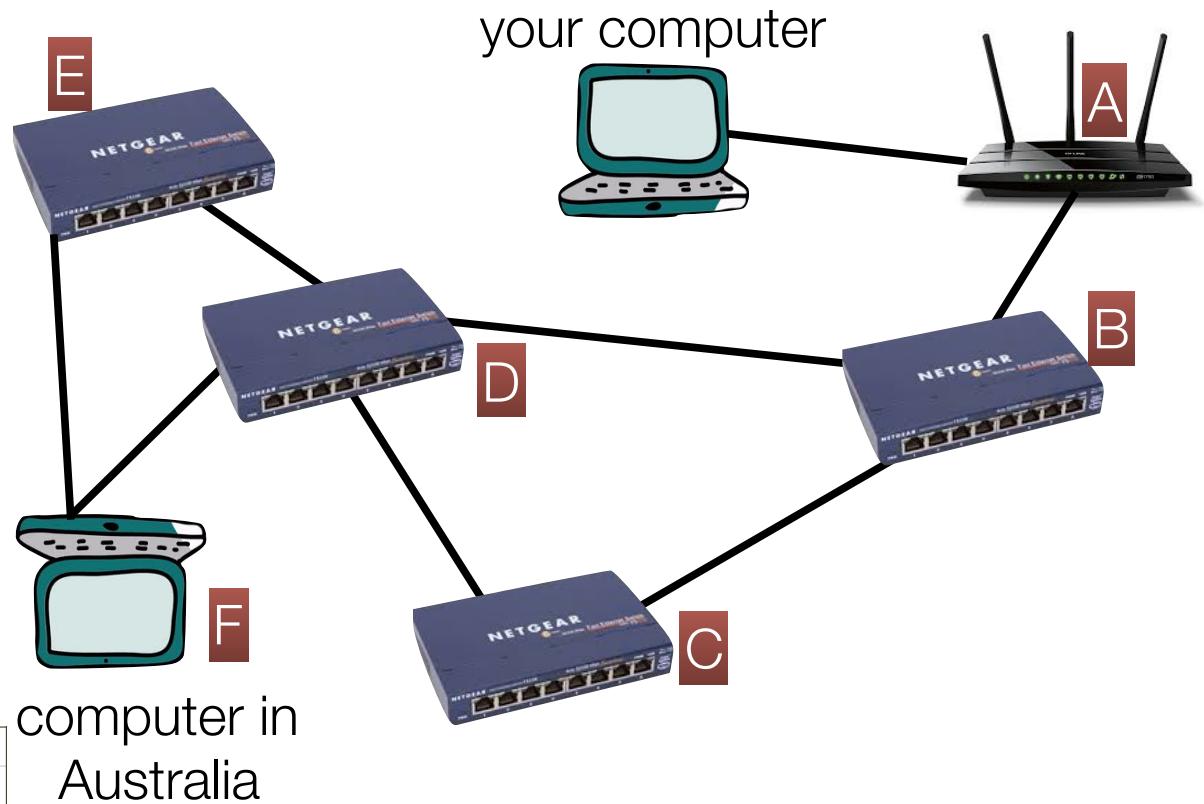
The destination computer has a name and an IP address, like this:

www.engineering.unsw.edu.au
IP address: 149.171.158.109

The first number denotes the "network address" and routers continually pass around information about how many "hops" they think it will take for them to get to all the networks.

E.g., for router C:

router	hops
A	2
B	1
C	-
D	1
E	2
F	2



The Internet: Let's simplify a bit

Each router knows its neighbors, and it has a copy of its neighbors' tables. So, B would have the following tables:

router	hops
A	-
B	1
C	3
D	2
E	3
F	3

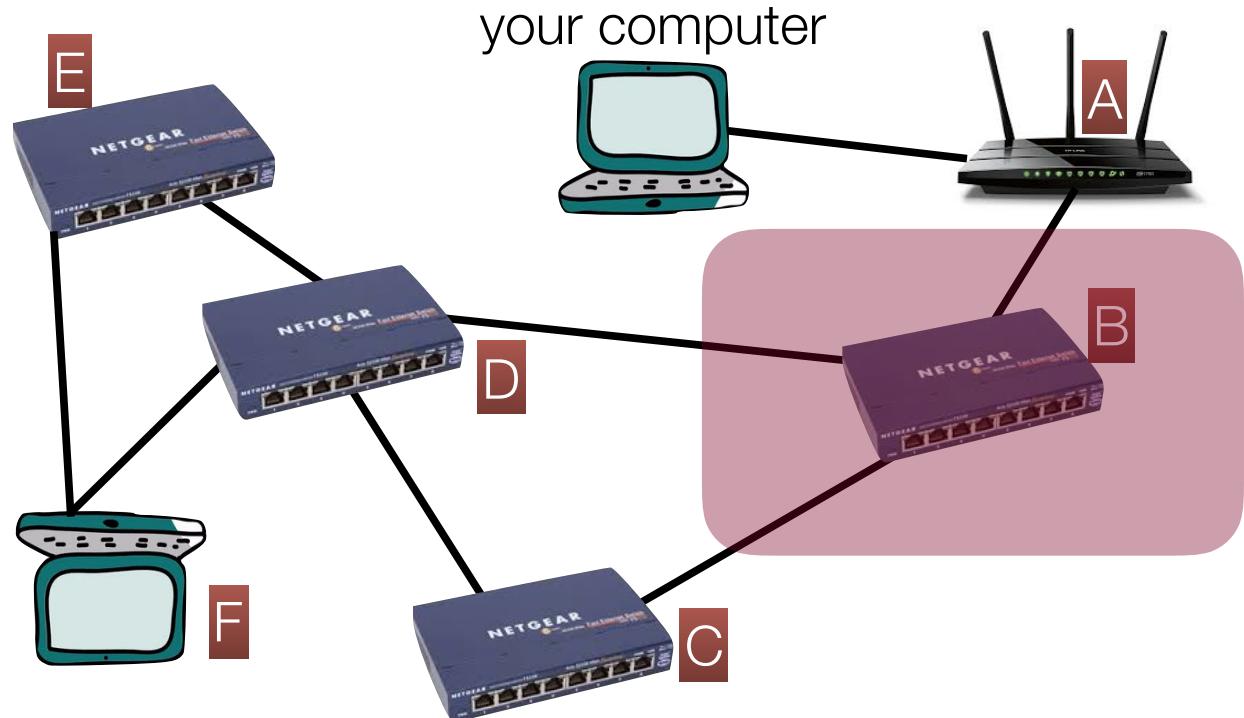
A

router	hops
A	2
B	1
C	-
D	1
E	2
F	2

C

router	hops
A	2
B	1
C	1
D	-
E	1
F	1

D



The Internet: Let's simplify a bit

If B wants to connect to F, it connects through its neighbor that reports the shortest path to F. Which router would it choose?

router	hops
A	-
B	1
C	3
D	2
E	3
F	3

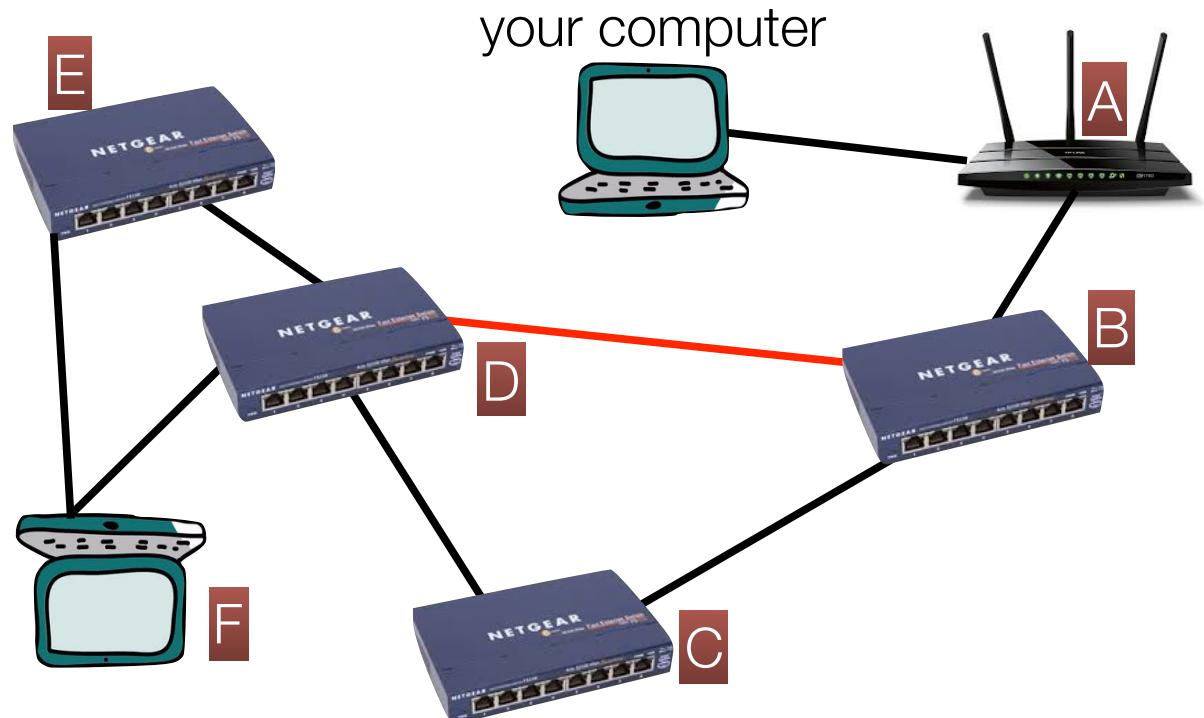
A

router	hops
A	2
B	1
C	-
D	1
E	2
F	2

C

D

router	hops
A	2
B	1
C	1
D	-
E	1
F	1



The Internet: Let's simplify a bit

If B wants to connect to F, it connects through its neighbor that reports the shortest path to F. Which router would it choose? D.

router	hops
A	-
B	1
C	3
D	2
E	3
F	3

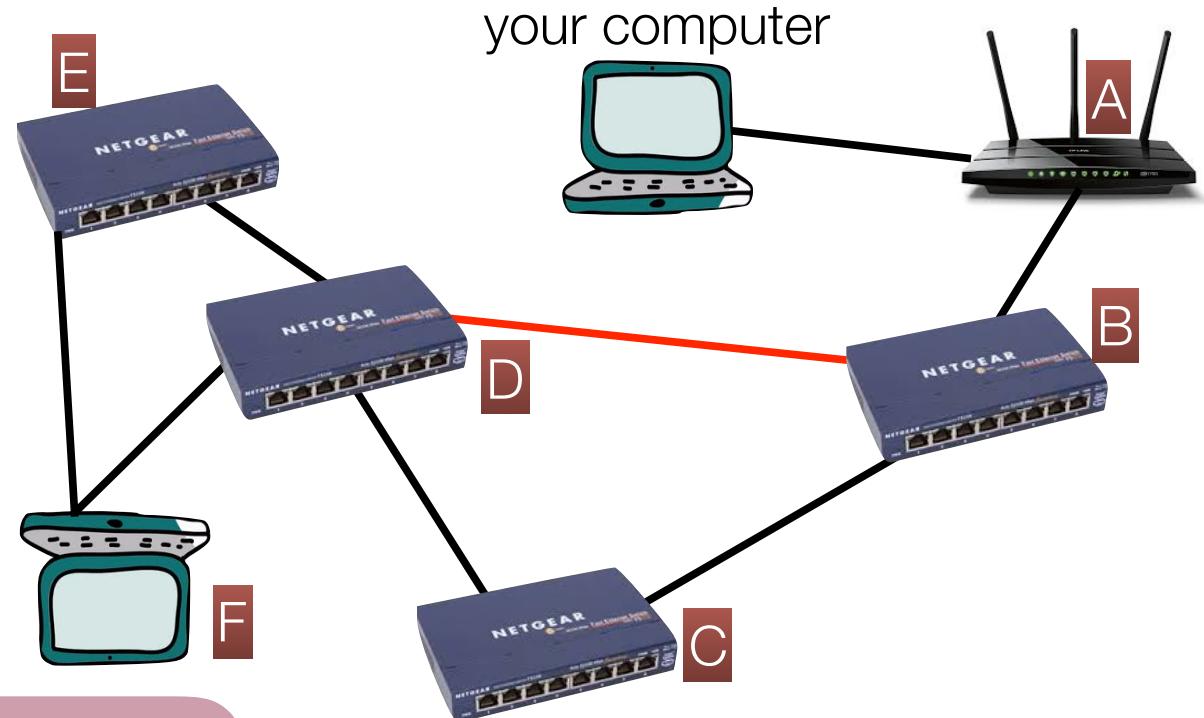
A

router	hops
A	2
B	1
C	-
D	1
E	2
F	2

C

D

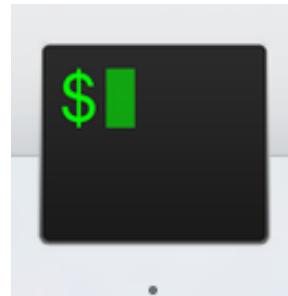
router	hops
A	2
B	1
C	1
D	-
E	1
F	1



Traceroute

We can use a program called "traceroute" to tell us the path between our computer and a different computer:

```
traceroute -I -e www.engineering.unsw.edu.au
```



Traceroute: Stanford Hops

```
traceroute -I -e www.engineering.unsw.edu.au
traceroute to www.engineering.unsw.edu.au (149.171.158.109), 64 hops max, 72 byte packets
 1 csmx-west-rtr.sunet (171.67.64.2)  7.414 ms  9.155 ms  8.288 ms
 2 gnat-2.sunet (172.24.70.12)  0.339 ms  1.532 ms  0.423 ms
 3 csmx-west-rtr-vl13866.sunet (171.64.66.2)  38.916 ms  10.506 ms  8.402 ms
 4 dca-rtr-vlan8.sunet (171.64.255.204)  0.530 ms  0.521 ms  0.713 ms
 5 dc-svl-agg4--stanford-10ge.cenic.net (137.164.50.157)  1.554 ms  1.653 ms  2.828 ms
 6 hpr-svl-hpr2--svl-agg4-10ge.cenic.net (137.164.26.249)  1.212 ms  1.161 ms  1.204 ms
 7 aarnet-2-is-jmb-778.sttlwa.pacificwave.net (207.231.245.4)  17.994 ms  17.998 ms  18.319 ms
 8 et-2-0-0.pe2.brwy.nsw.aarnet.net.au (113.197.15.98)  160.020 ms  160.234 ms  159.922 ms
 9 et-3-3-0.pe1.brwy.nsw.aarnet.net.au (113.197.15.148)  160.285 ms  160.076 ms  160.118 ms
10 138.44.5.1 (138.44.5.1)  160.124 ms  160.138 ms  160.068 ms
11 ombcr1-te-1-5.gw.unsw.edu.au (149.171.255.106)  160.090 ms  160.381 ms  160.185 ms
12 r1dcdnex1-po-2.gw.unsw.edu.au (149.171.255.178)  160.909 ms  160.847 ms  160.921 ms
13 dcfwl-ae-1-3049.gw.unsw.edu.au (129.94.254.60)  160.592 ms  160.558 ms  160.949 ms
14 www.engineering.unsw.edu.au (149.171.158.109)  160.978 ms  161.184 ms  160.987 ms
```



Traceroute: CENIC

```
traceroute -I -e www.engineering.unsw.edu.au
traceroute to www.engineering.unsw.edu.au (149.171.158.109), 64 hops max, 72 byte packets
 1 csmx-west-rtr.sunet (171.67.64.2)  7.414 ms  9.155 ms  8.288 ms
 2 gnat-2.sunet (172.24.70.12)  0.339 ms  1.532 ms  0.423 ms
 3 csmx-west-rtr-v13866.sunet (171.64.66.2)  38.916 ms  10.506 ms  8.402 ms
 4 dca-rtr-vlan8.sunet (171.64.255.204)  0.530 ms  0.521 ms  0.713 ms
 5 dc-svl-agg4--stanford-10ge.cenic.net (137.164.50.157)  1.554 ms  1.653 ms  2.828 ms
 6 hpr-svl-hpr2--svl-agg4-10ge.cenic.net (137.164.26.249)  1.212 ms  1.161 ms  1.204 ms
 7 aarnet-2-is-jmb-778.sttlwa.pacificwave.net (207.231.245.4)  17.994 ms  17.998 ms  18.319 ms
 8 et-2-0-0.pe2.brwy.nsw.aarnet.net.au (113.197.15.98)  160.020 ms  160.234 ms  159.922 ms
 9 et-3-3-0.pe1.brwy.nsw.aarnet.net.au (113.197.15.148)  160.285 ms  160.076 ms  160.118 ms
10 138.44.5.1 (138.44.5.1)  160.124 ms  160.138 ms  160.068 ms
11 ombcr1-te-1-5.gw.unsw.edu.au (149.171.255.106)  160.090 ms  160.381 ms  160.185 ms
12 r1dcdnex1-po-2.gw.unsw.edu.au (149.171.255.178)  160.909 ms  160.847 ms  160.921 ms
13 dcfwl-ae-1-3049.gw.unsw.edu.au (129.94.254.60)  160.592 ms  160.558 ms  160.949 ms
14 www.engineering.unsw.edu.au (149.171.158.109)  160.978 ms  161.184 ms  160.987 ms
```

The Corporation for Education Network Initiatives in California (**CENIC**) is a nonprofit corporation formed in 1996 to provide high-performance, high-bandwidth networking services to [California](#) universities and research institutions (source: Wikipedia)



Traceroute: Pacificwave (Seattle)

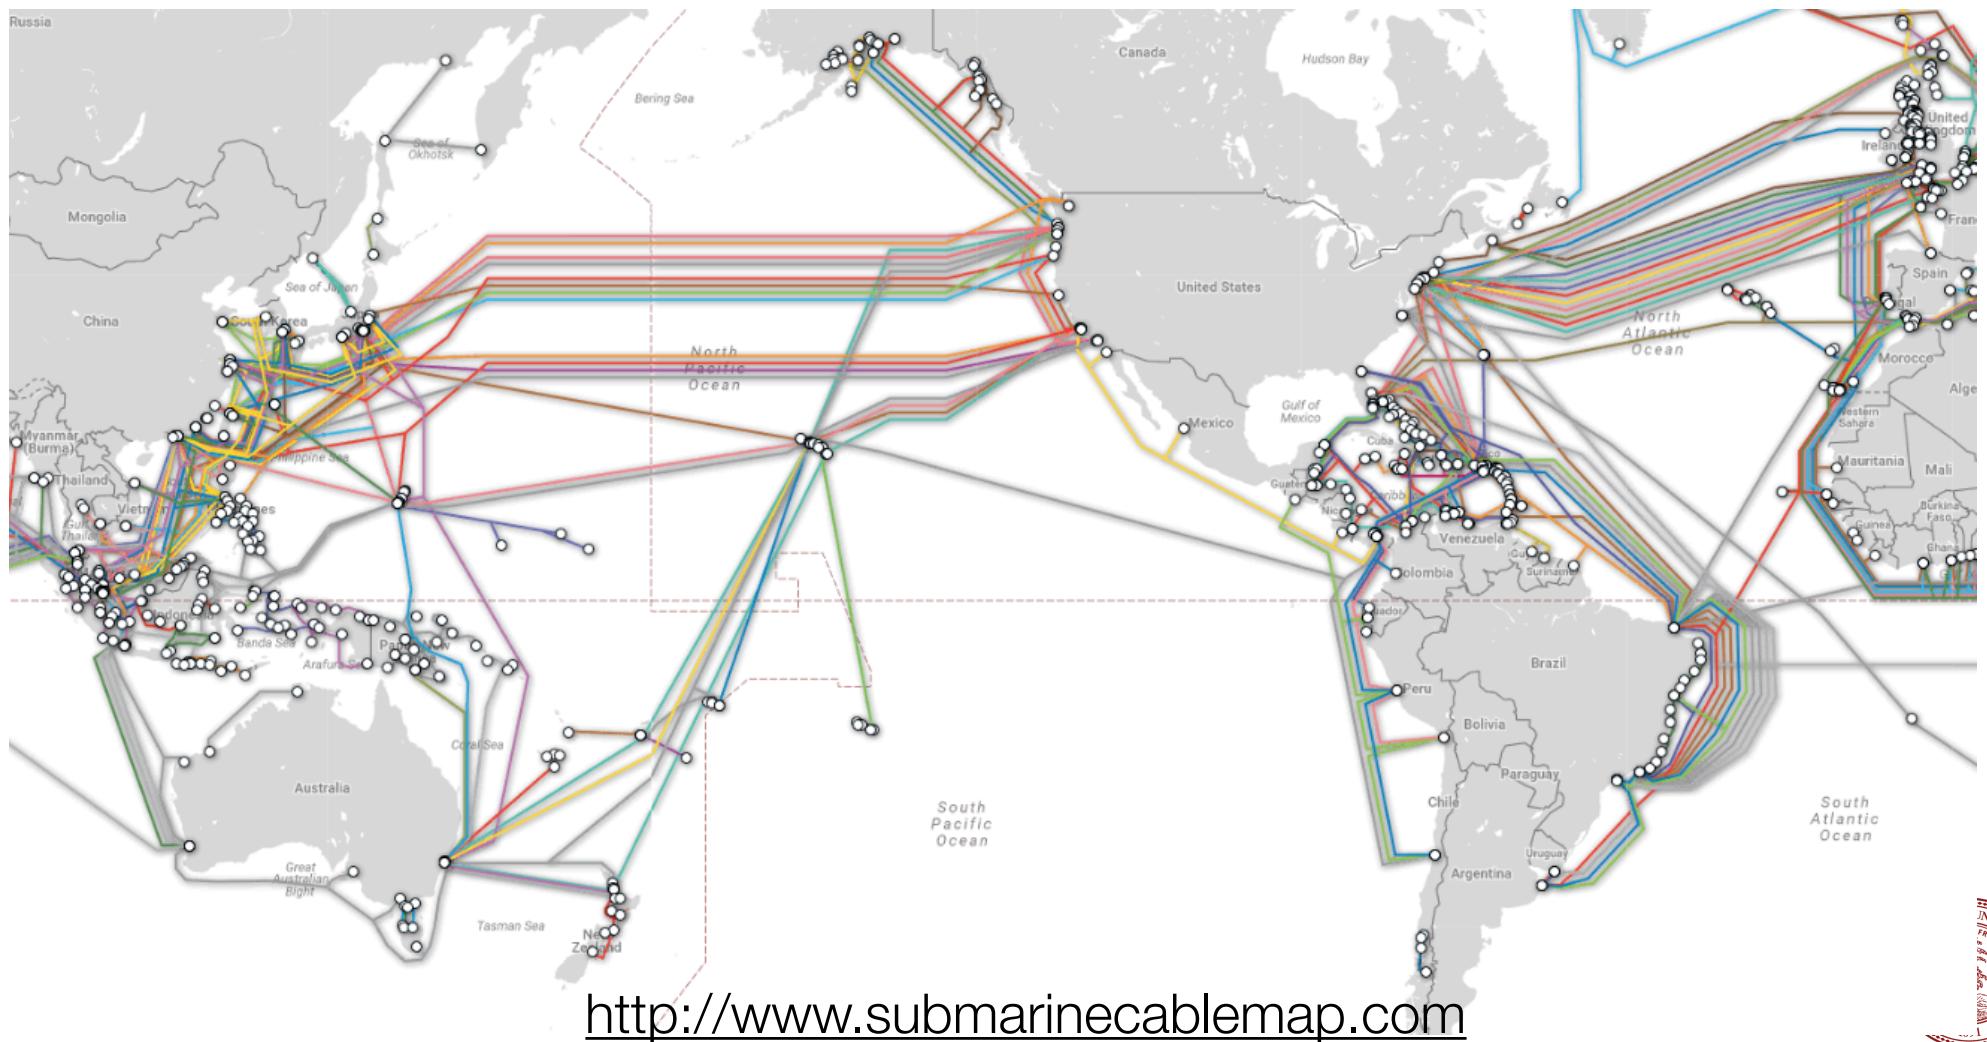
```
traceroute -I -e www.engineering.unsw.edu.au
traceroute to www.engineering.unsw.edu.au (149.171.158.109), 64 hops max, 72 byte packets
 1 csmx-west-rtr.sunet (171.67.64.2)  7.414 ms  9.155 ms  8.288 ms
 2 gnat-2.sunet (172.24.70.12)  0.339 ms  1.532 ms  0.423 ms
 3 csmx-west-rtr-v13866.sunet (171.64.66.2)  38.916 ms  10.506 ms  8.402 ms
 4 dca-rtr-vlan8.sunet (171.64.255.204)  0.530 ms  0.521 ms  0.713 ms
 5 dc-svl-agg4--stanford-10ge.cenic.net (137.164.50.157)  1.554 ms  1.653 ms  2.828 ms
 6 hpr-svl-hpr2--svl-agg4-10ge.cenic.net (137.164.26.249)  1.212 ms  1.161 ms  1.204 ms
 7 aarnet-2-is-jmb-778.sttlwa.pacificwave.net (207.231.245.4)  17.994 ms  17.998 ms  18.319 ms
 8 et-2-0-0.pe2.brwy.nsw.aarnet.net.au (113.197.15.98)  160.020 ms  160.234 ms  159.922 ms
 9 et-3-3-0.pe1.brwy.nsw.aarnet.net.au (113.197.15.148)  160.285 ms  160.076 ms  160.118 ms
10 138.44.5.1 (138.44.5.1)  160.124 ms  160.138 ms  160.068 ms
                               gw.unsw.edu.au (149.171.255.106)  160.090 ms  160.381 ms  160.185 ms
                               .gw.unsw.edu.au (149.171.255.178)  160.909 ms  160.847 ms  160.921 ms
                               9.gw.unsw.edu.au (129.94.254.60)  160.592 ms  160.558 ms  160.949 ms
                               g.unsw.edu.au (149.171.158.109)  160.978 ms  161.184 ms  160.987 ms
```



Pass Internet traffic directly with other major national and international networks, including U.S. federal agencies and many Pacific Rim R&E networks (source: <http://www.pnwgp.net/services/pacific-wave-peering-exchange/>)



Traceroute: Oregon to Australia - underwater!



Traceroute: Australia

```
traceroute -I -e www.engineering.unsw.edu.au
traceroute to www.engineering.unsw.edu.au (149.171.158.109), 64 hops max, 72 byte packets
 1 csmx-west-rtr.sunet (171.67.64.2)  7.414 ms  9.155 ms  8.288 ms
 2 gnat-2.sunet (172.24.70.12)  0.339 ms  1.532 ms  0.423 ms
 3 csmx-west-rtr-v13866.sunet (171.64.66.2)  38.916 ms  10.506 ms  8.402 ms
 4 dca-rtr-vlan8.sunet (171.64.255.204)  0.530 ms  0.521 ms  0.713 ms
 5 dc-svl-agg4--stanford-10ge.cenic.net (137.164.50.157)  1.554 ms  1.653 ms  2.828 ms
 6 hpr-svl-hpr2--svl-agg4-10ge.cenic.net (137.164.26.249)  1.212 ms  1.161 ms  1.204 ms
 7 aarnet-2-is-jmb-778.sttlwa.pacificwave.net (207.231.245.4)  17.994 ms  17.998 ms  18.319 ms
 8 et-2-0-0.pe2.brwy.nsw.aarnet.net.au (113.197.15.98)  160.020 ms  160.234 ms  159.922 ms
 9 et-3-3-0.pe1.brwy.nsw.aarnet.net.au (113.197.15.148)  160.285 ms  160.076 ms  160.118 ms
10 138.44.5.1 (138.44.5.1)  160.124 ms  160.138 ms  160.068 ms
11 ombcrl-te-1-5.gw.unsw.edu.au (149.171.255.106)  160.090 ms  160.381 ms  160.185 ms
12 r1dcdnex1-po-2.gw.unsw.edu.au (149.171.255.178)  160.909 ms  160.847 ms  160.921 ms
13 dcfwl-ae-1-3049.gw.unsw.edu.au (129.94.254.60)  160.592 ms  160.558 ms  160.949 ms
14 www.engineering.unsw.edu.au (149.171.158.109)  160.978 ms  161.184 ms  160.987 ms
```



Traceroute: University of New South Wales

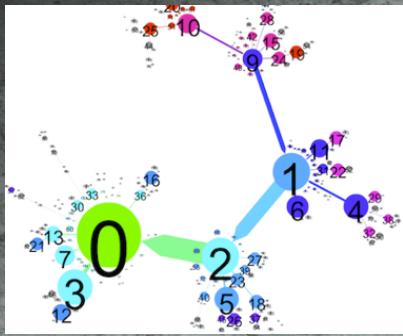
```
traceroute -I -e www.engineering.unsw.edu.au
traceroute to www.engineering.unsw.edu.au (149.171.158.109), 64 hops max, 72 byte packets
 1 csmx-west-rtr.sunet (171.67.64.2)  7.414 ms  9.155 ms  8.288 ms
 2 gnat-2.sunet (172.24.70.12)  0.339 ms  1.532 ms  0.423 ms
 3 csmx-west-rtr-v13866.sunet (171.64.66.2)  38.916 ms  10.506 ms  8.402 ms
 4 dca-rtr-vlan8.sunet (171.64.255.204)  0.530 ms  0.521 ms  0.713 ms
 5 dc-svl-agg4--stanford-10ge.cenic.net (137.164.50.157)  1.554 ms  1.653 ms  2.828 ms
 6 hpr-svl-hpr2--svl-agg4-10ge.cenic.net (137.164.26.249)  1.212 ms  1.161 ms  1.204 ms
 7 aarnet-2-is-jmb-778.sttlwa.pacificwave.net (207.231.245.4)  17.994 ms  17.998 ms  18.319 ms
 8 et-2-0-0.pe2.brwy.nsw.aarnet.net.au (113.197.15.98)  160.020 ms  160.234 ms  159.922 ms
 9 et-3-3-0.pe1.brwy.nsw.aarnet.net.au (113.197.15.148)  160.285 ms  160.076 ms  160.118 ms
10 138.44.5.1 (138.44.5.1)  160.124 ms  160.138 ms  160.068 ms
11 ombcr1-te-1-5.gw.unsw.edu.au (149.171.255.106)  160.090 ms  160.381 ms  160.185 ms
12 r1dcdnex1-po-2.gw.unsw.edu.au (149.171.255.178)  160.909 ms  160.847 ms  160.921 ms
13 dcfwl-ae-1-3049.gw.unsw.edu.au (129.94.254.60)  160.592 ms  160.558 ms  160.949 ms
14 www.engineering.unsw.edu.au (149.171.158.109)  160.978 ms  161.184 ms  160.987 ms
```

161 milliseconds to get to the final computer



More Real Graphs: Chris Piech's Research!

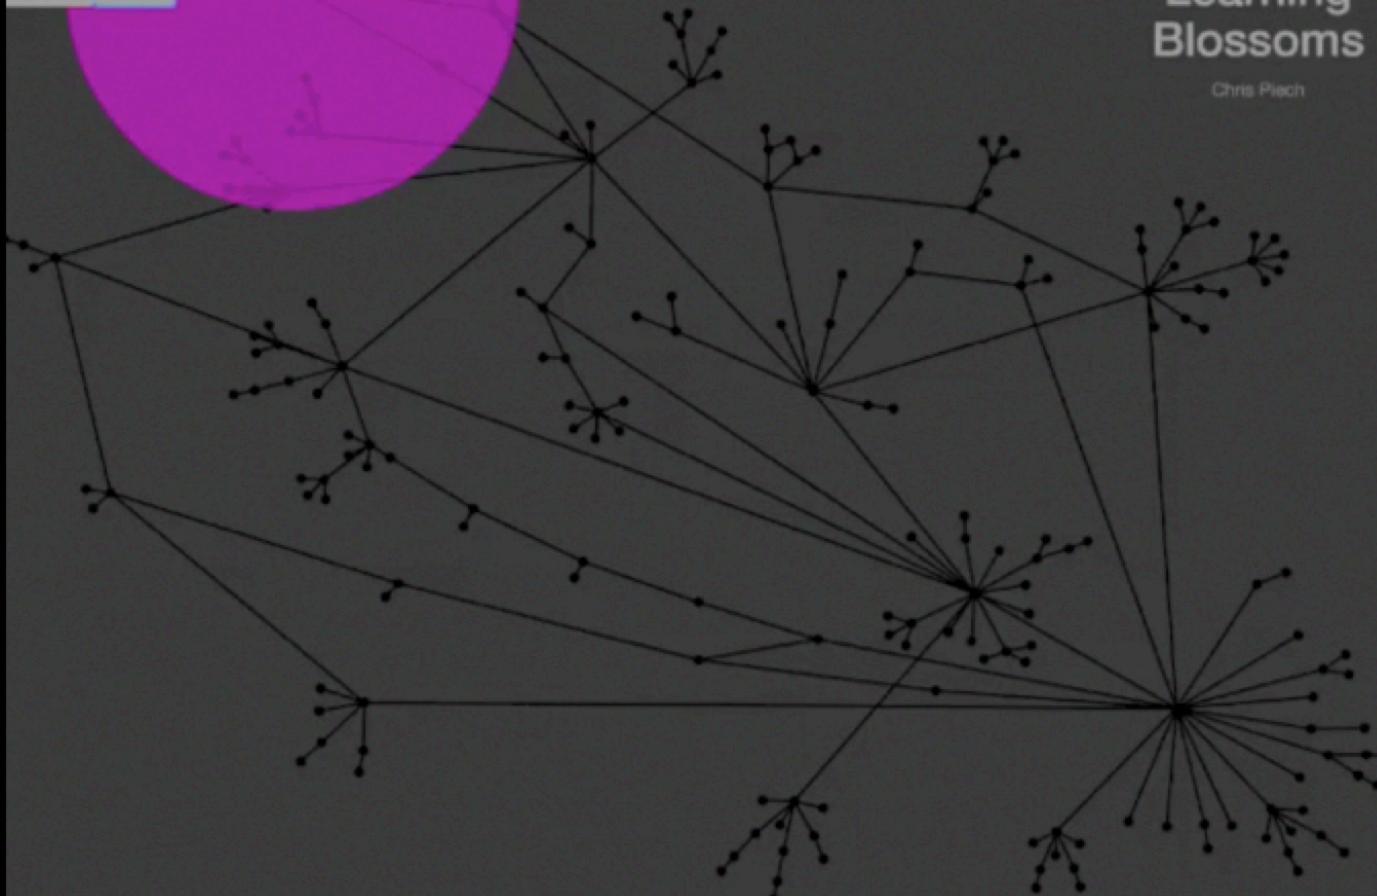
Autonomously Generating Hints by Inferring Problem Solving Policies

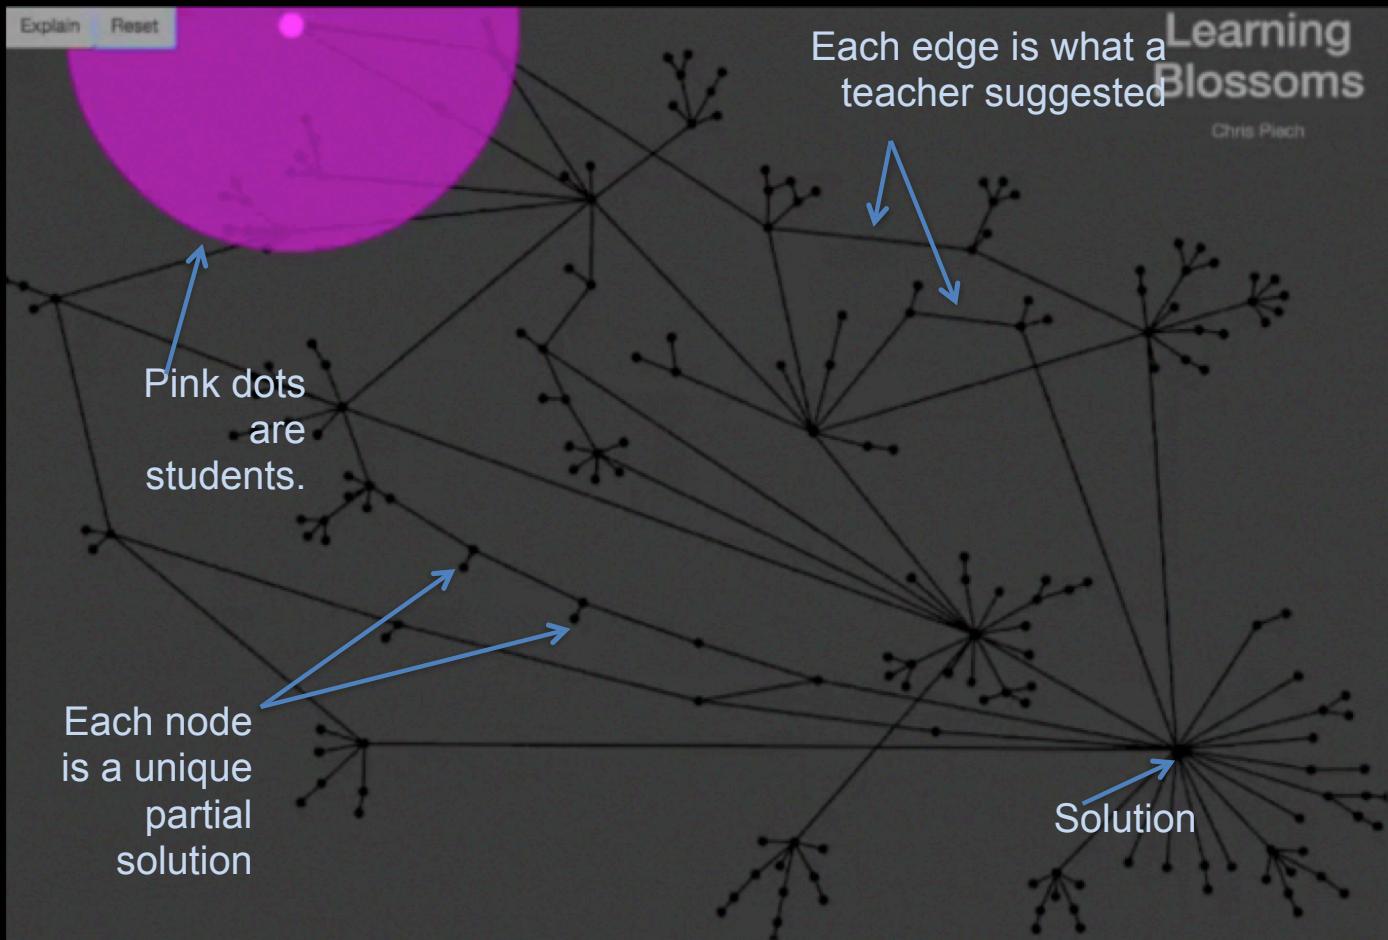


Explain Reset

Learning Blossoms

Chris Piesch

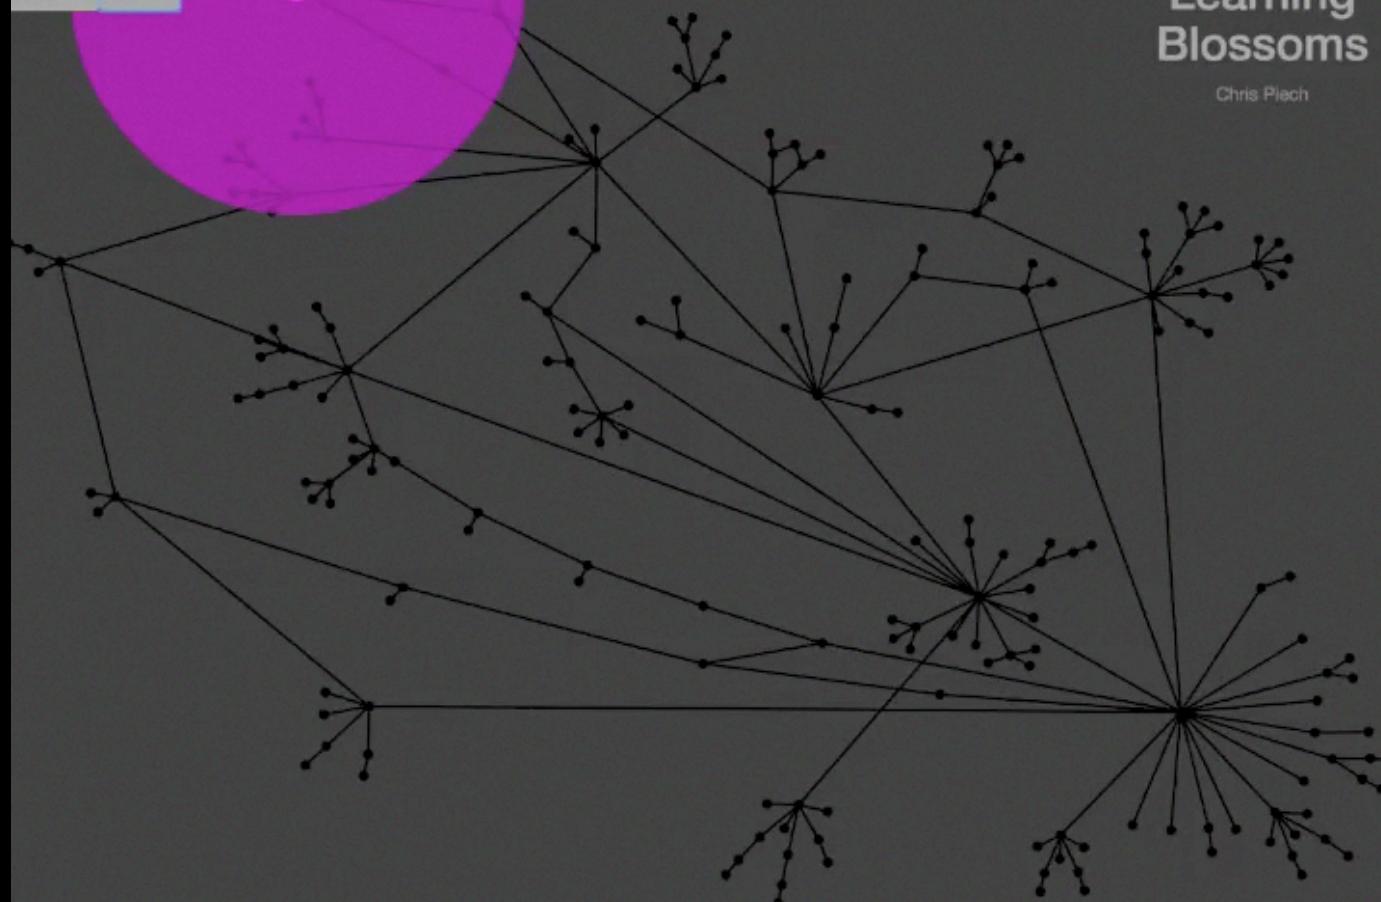




Explain Reset

Learning Blossoms

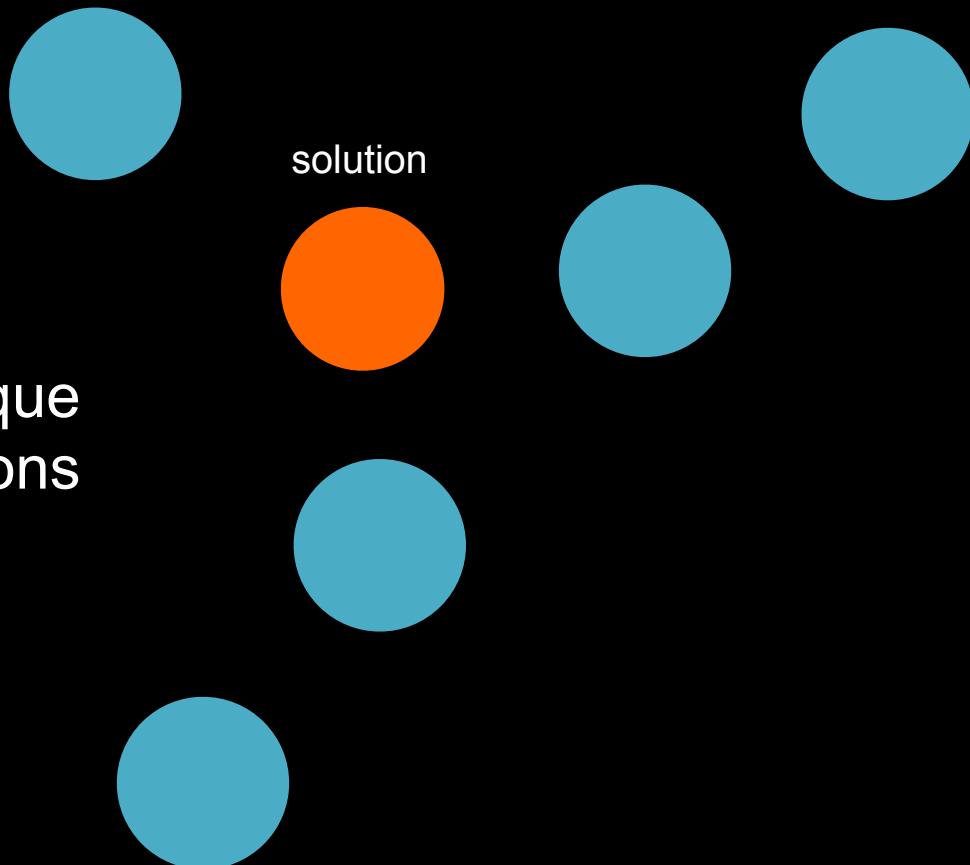
Chris Piech



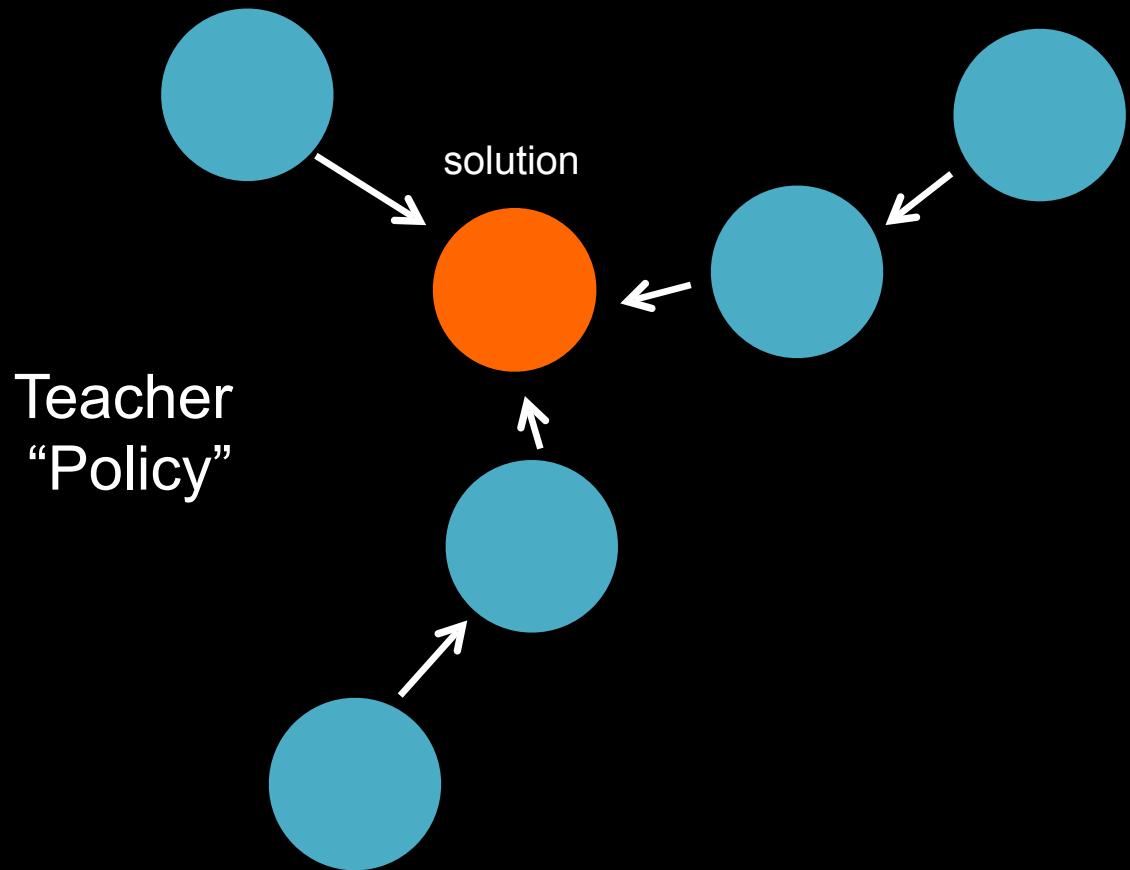
Machine Learning Problem

Unique
submissions

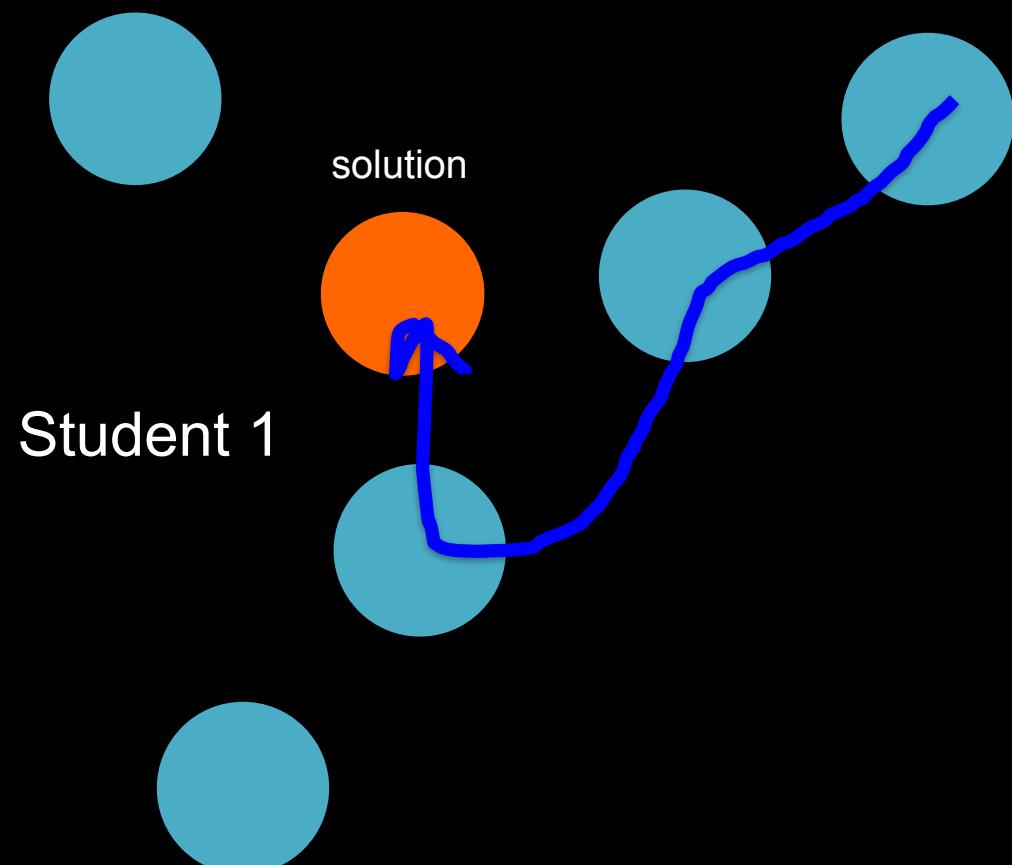
solution



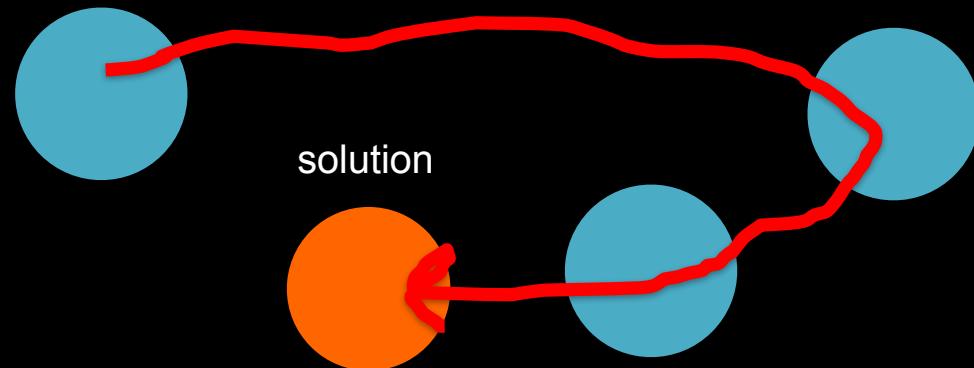
Machine Learning Problem



Example Student



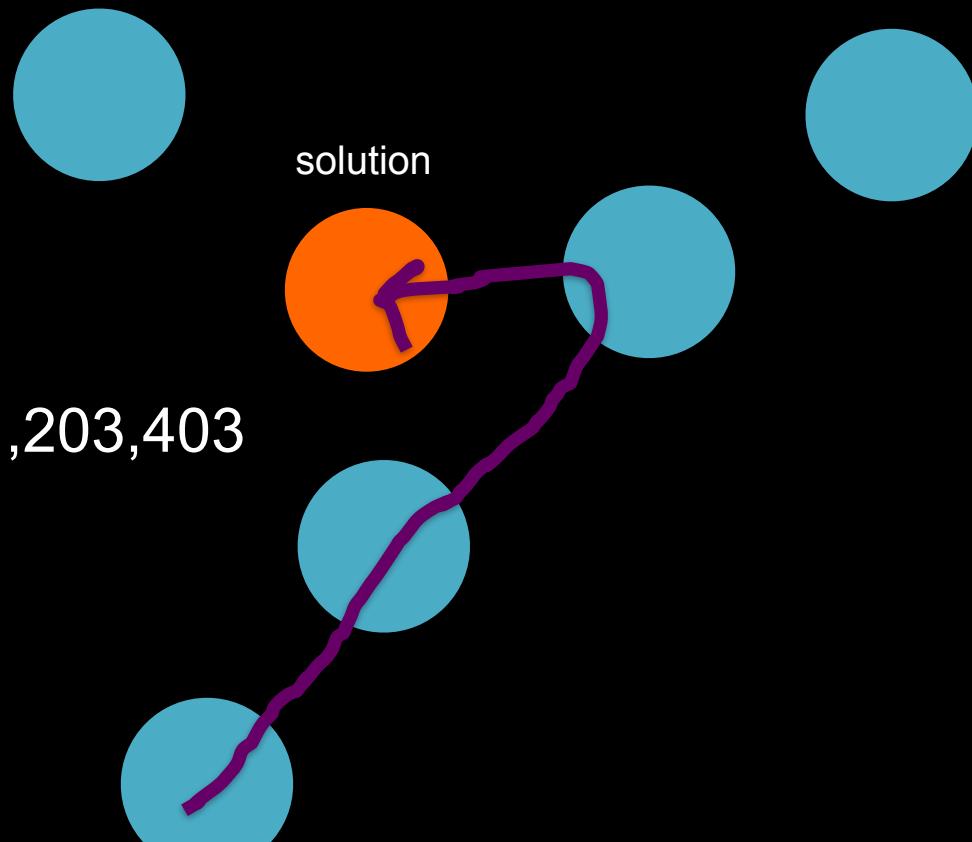
Example Student



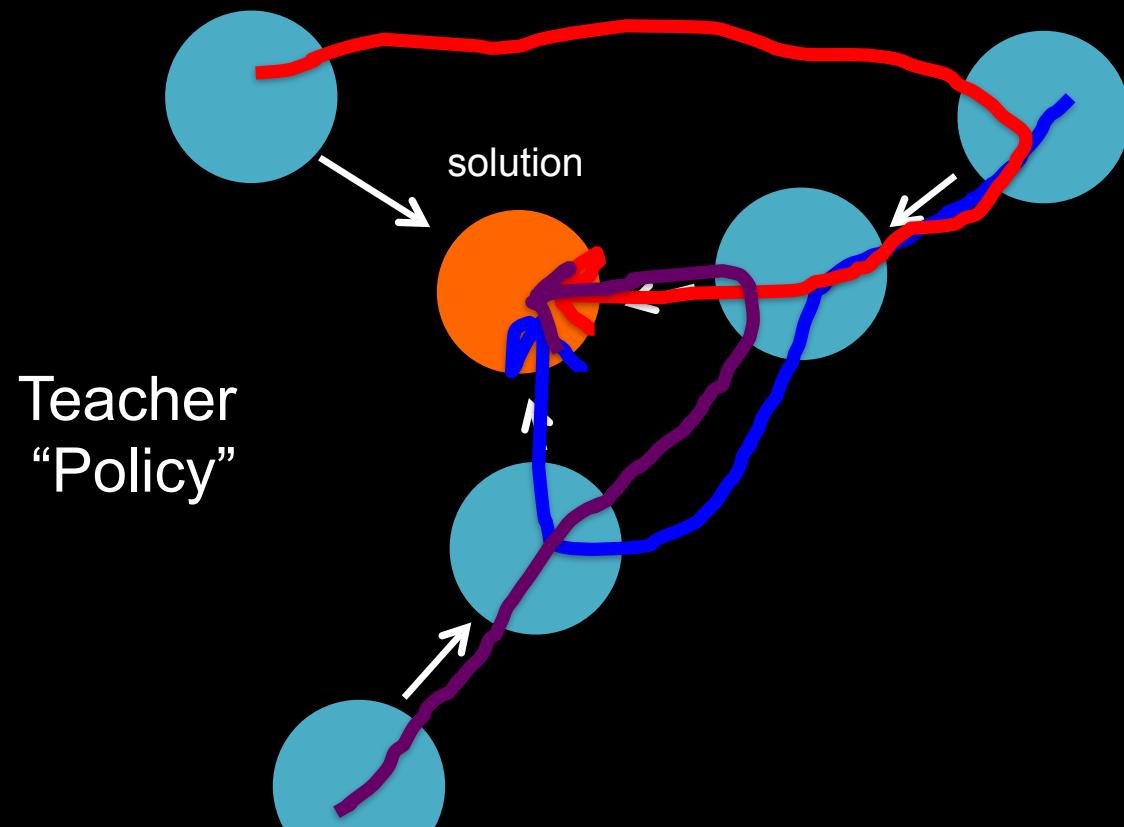
Student 2

Example Student

Student 1,203,403



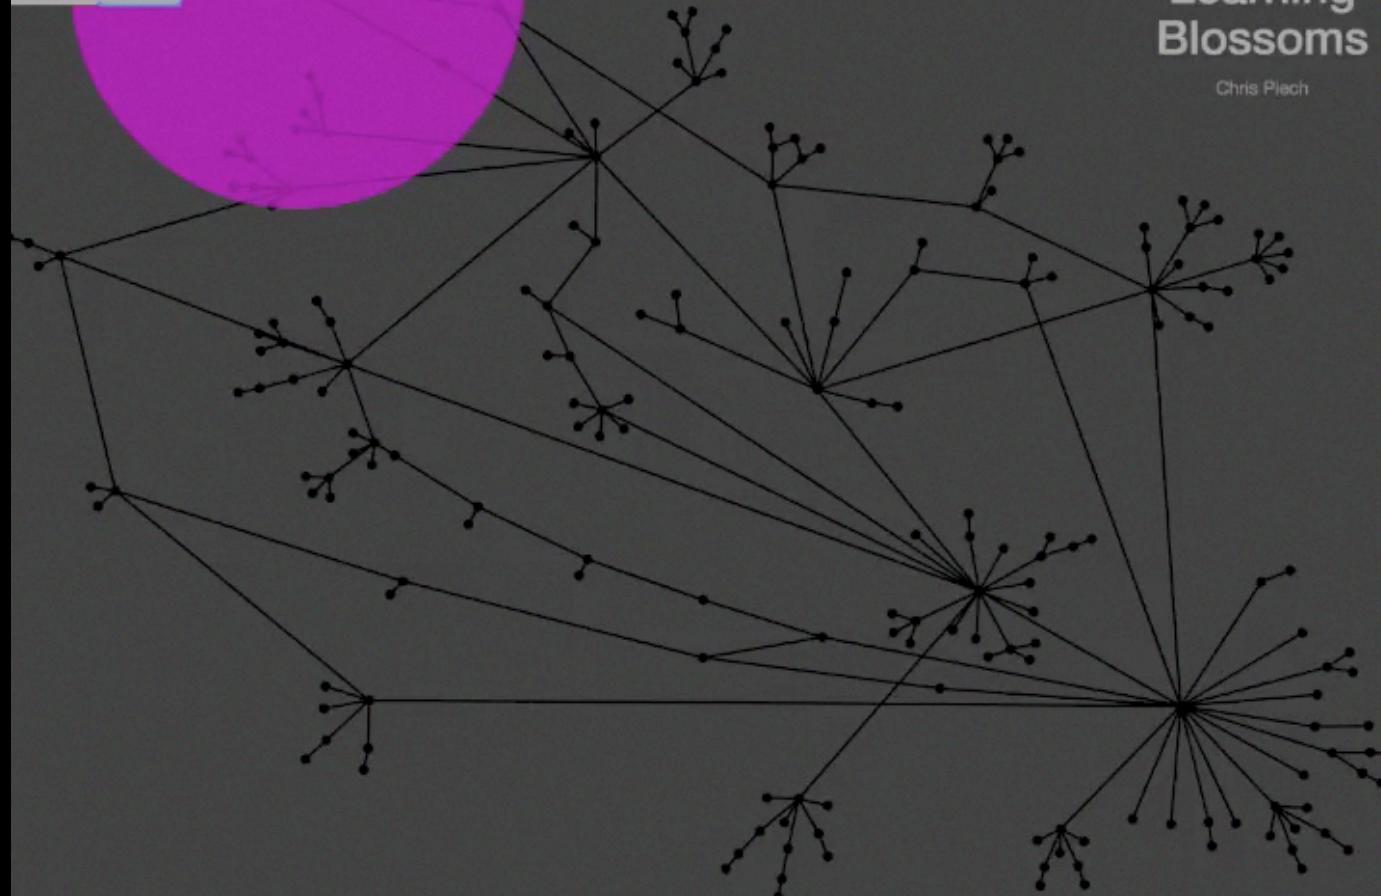
Machine Learning Problem



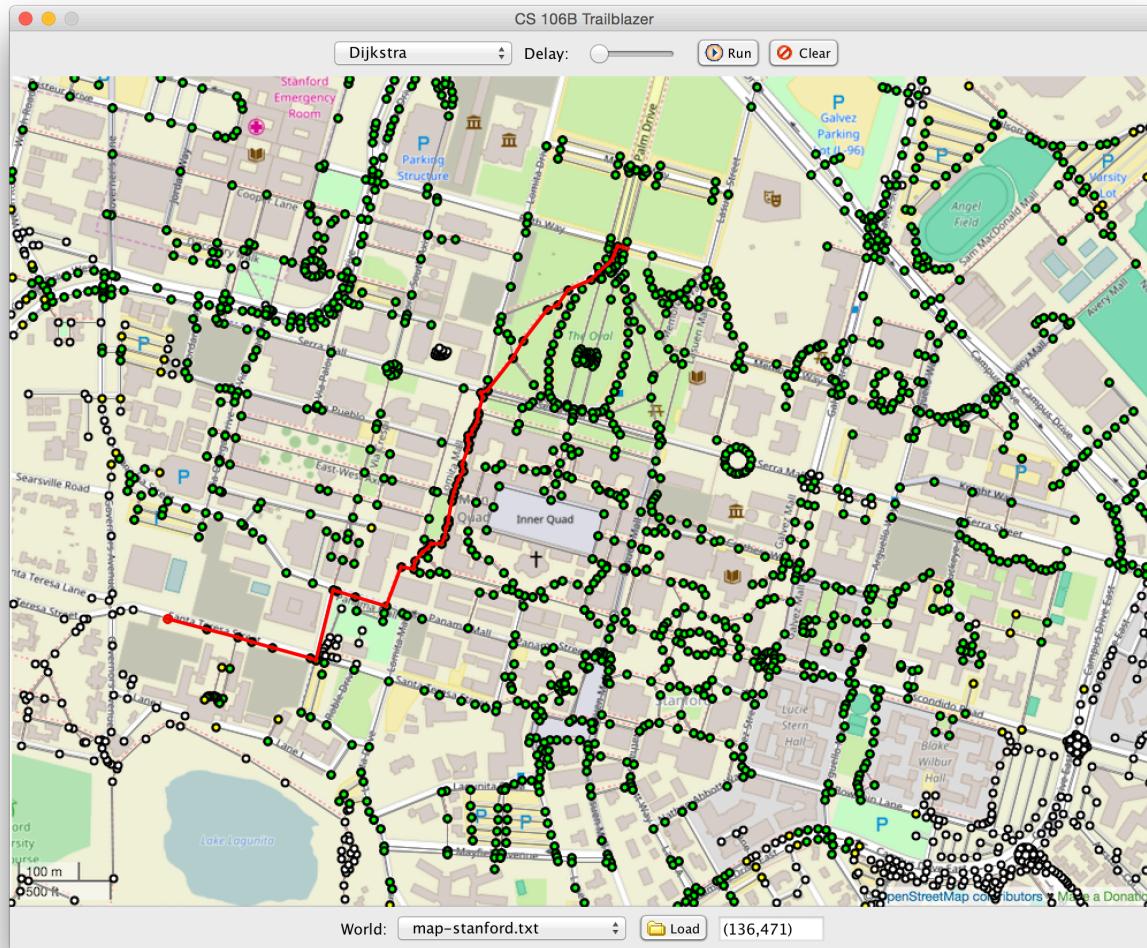
Explain Reset

Learning Blossoms

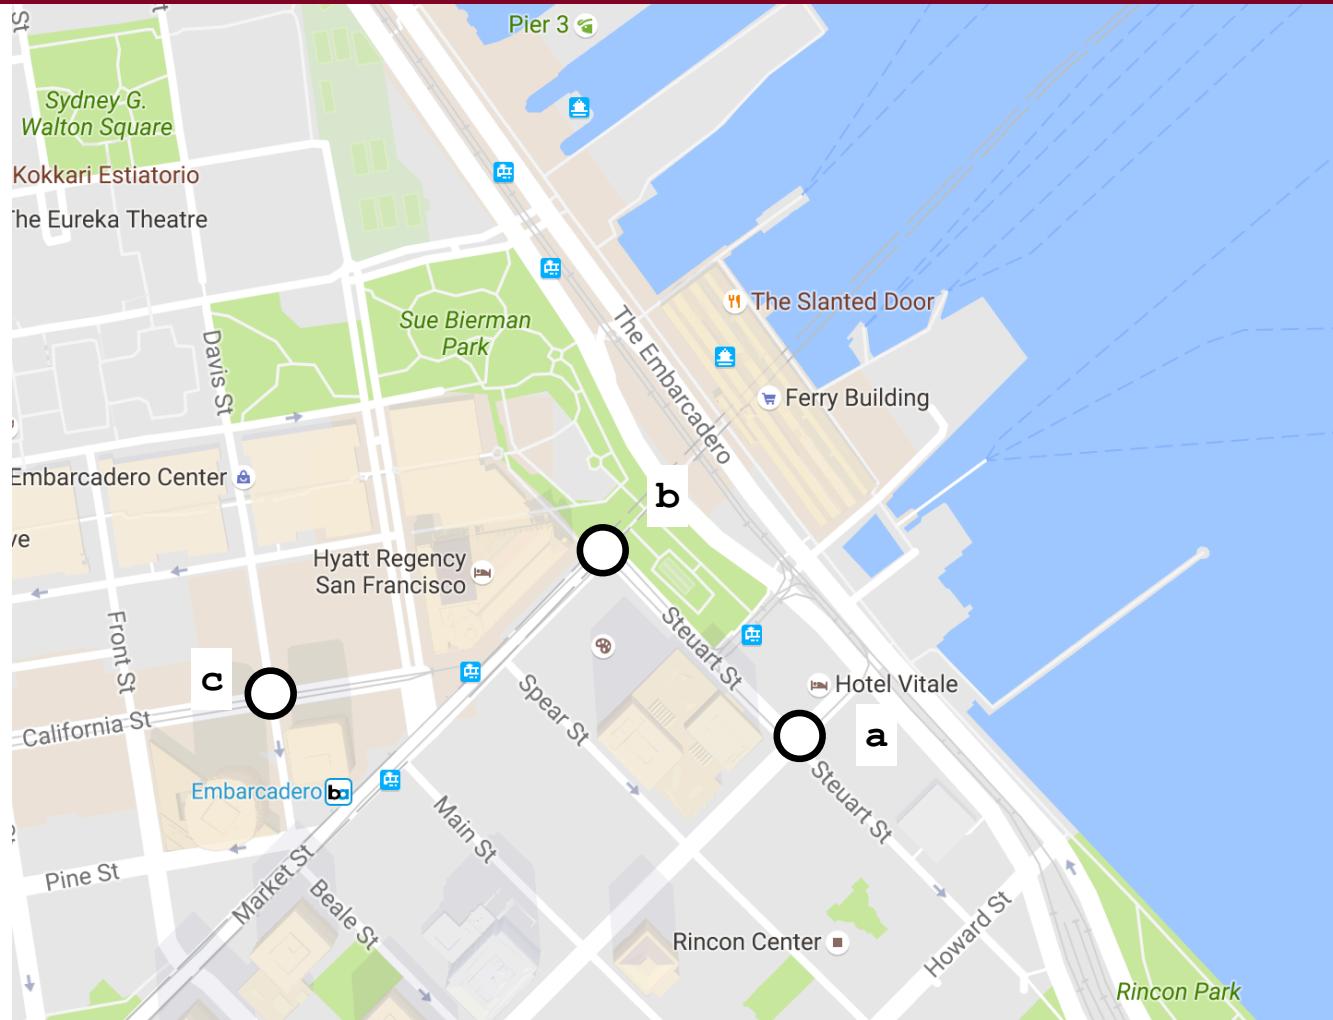
Chris Piech



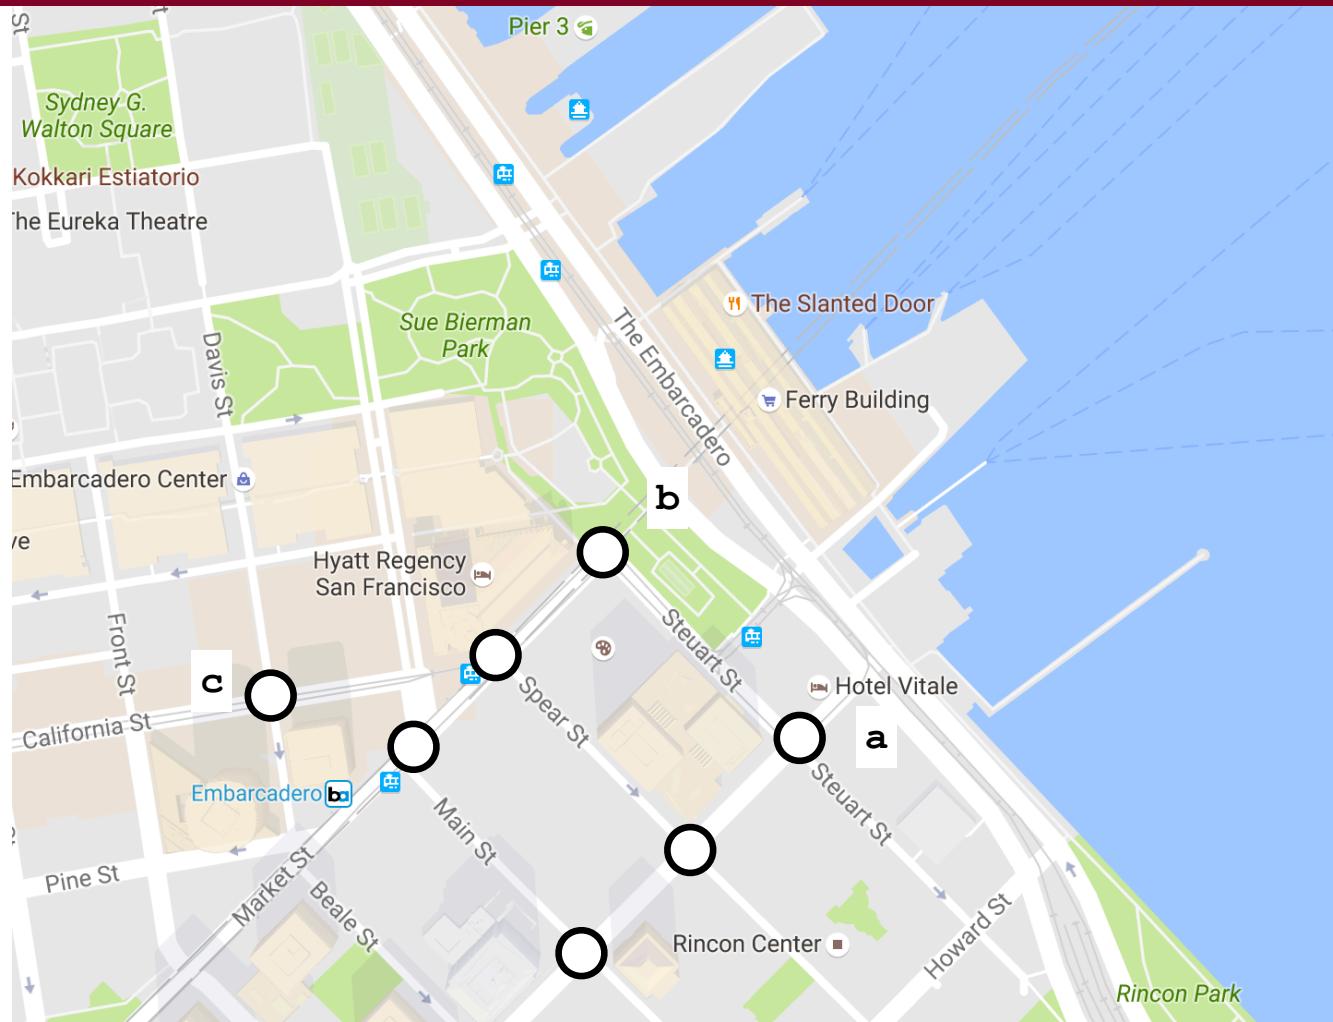
Back to Trailblazer



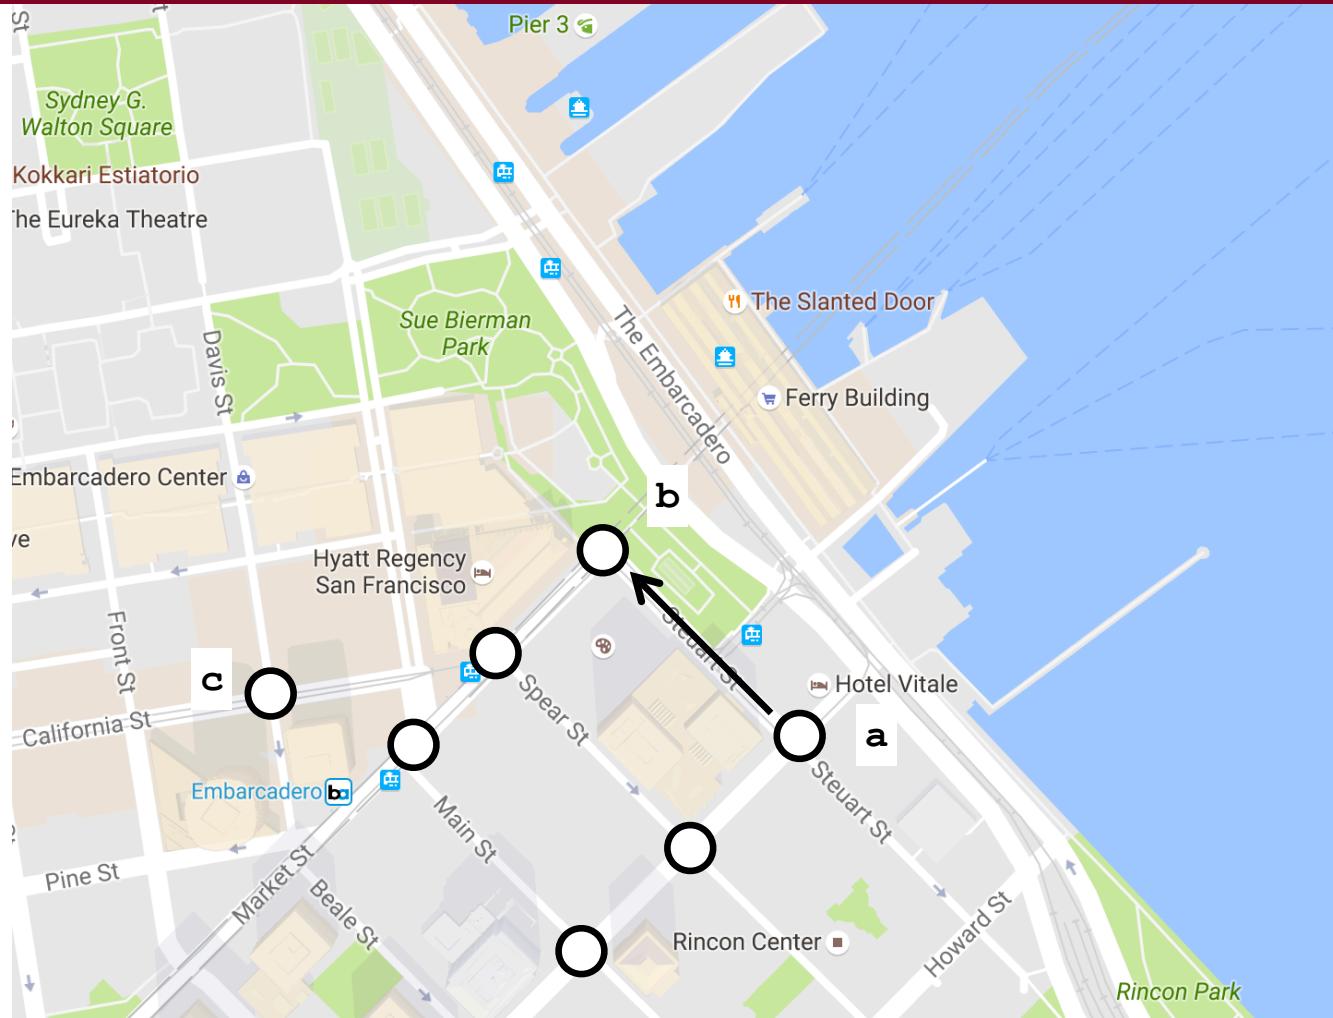
Road Map Node



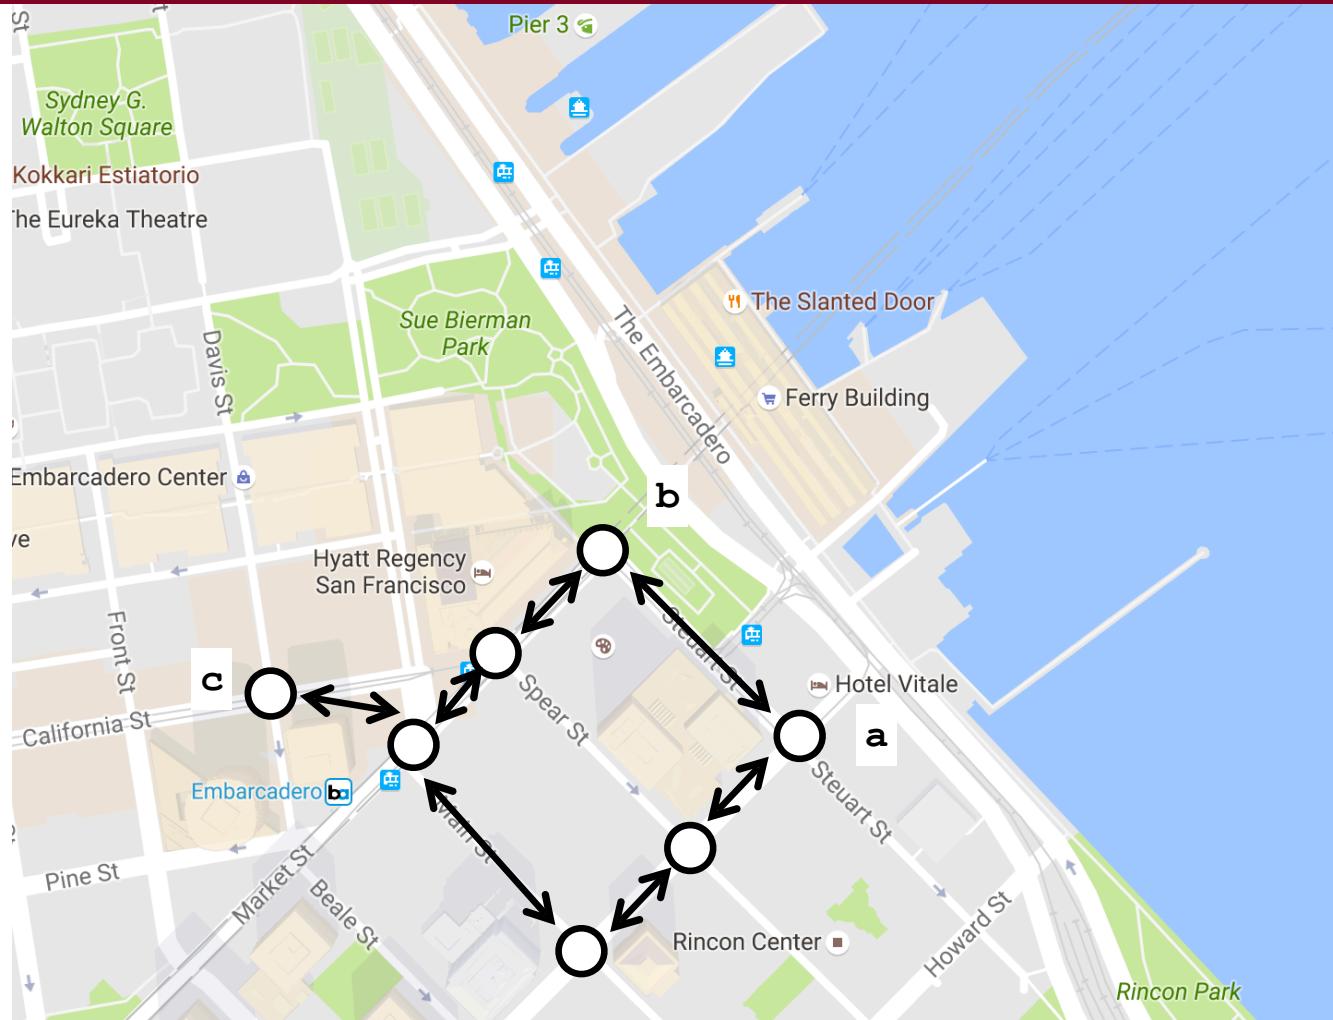
Road Map Node



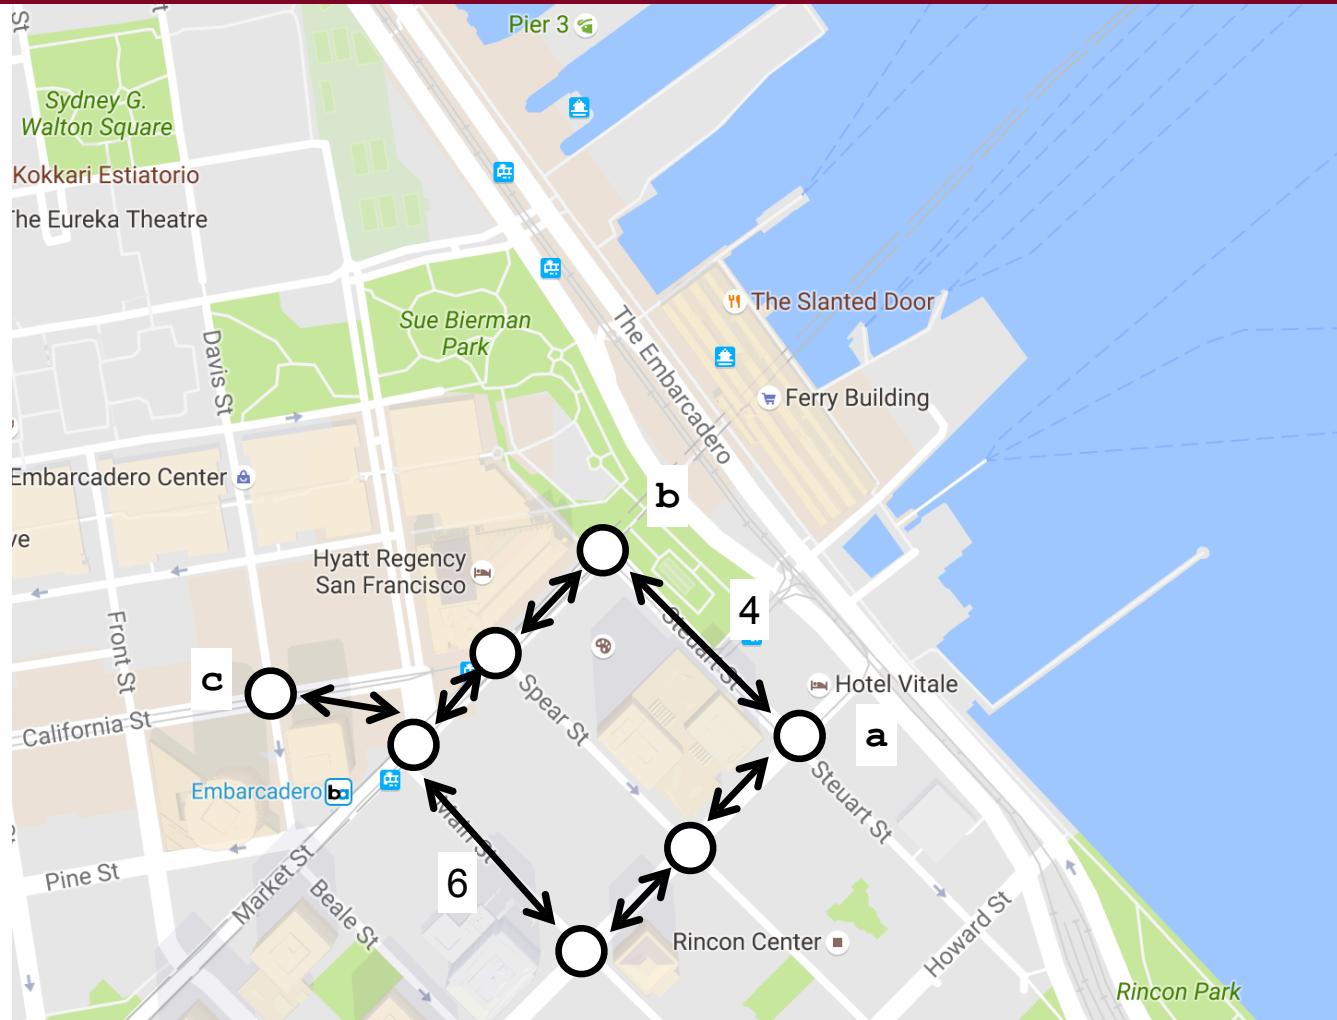
Road Map Edge



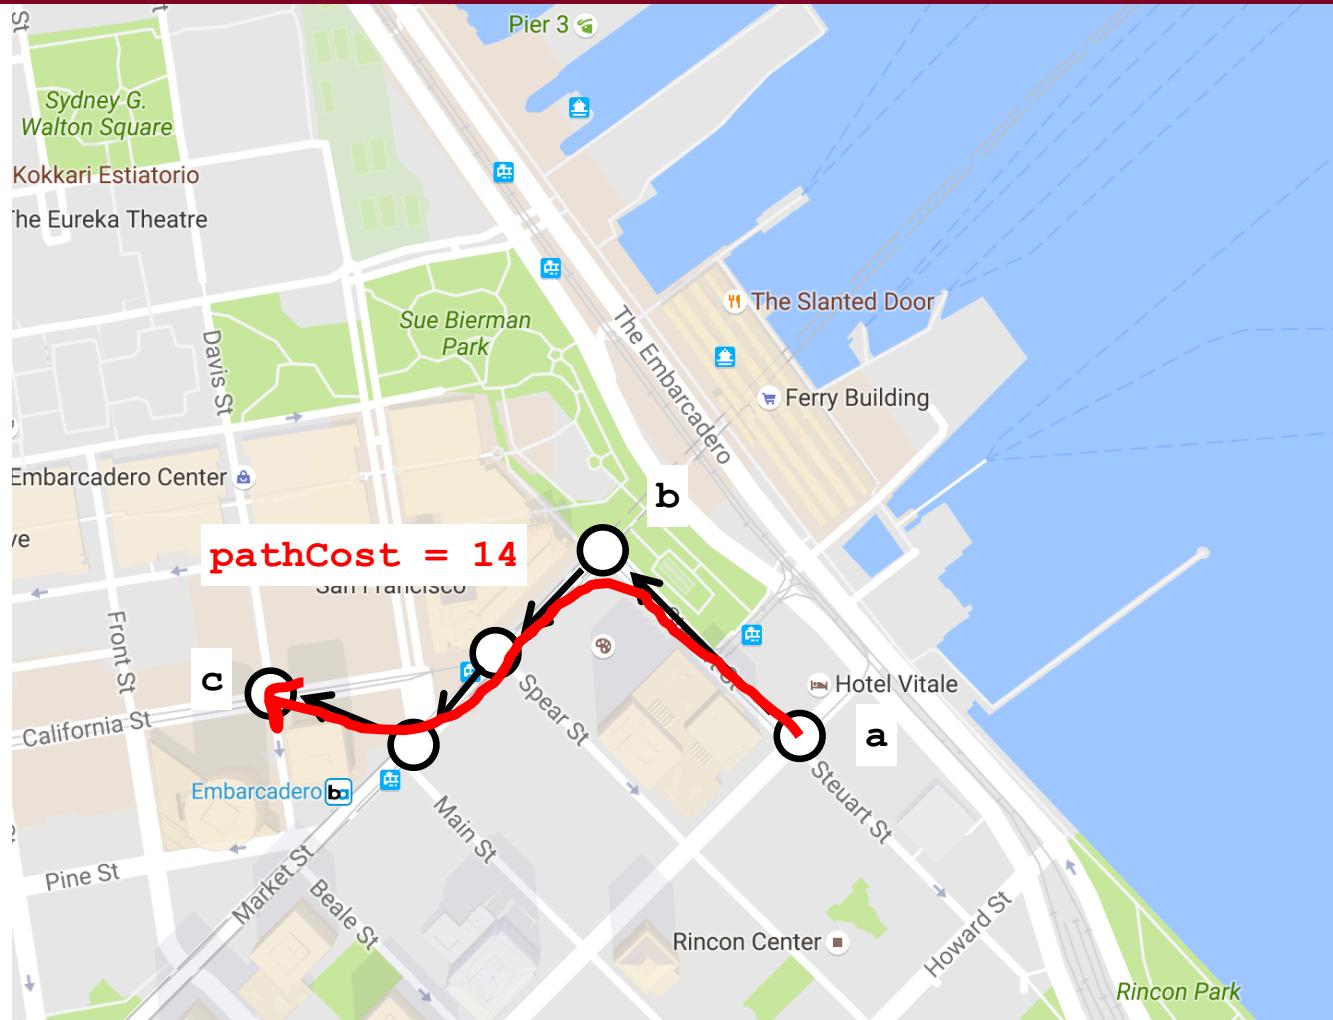
Road Map Edge



Road Map Edge Cost



Road Map Path Cost



Could Google Just Precompute?

How many nodes in google maps graph?

~ 75 million

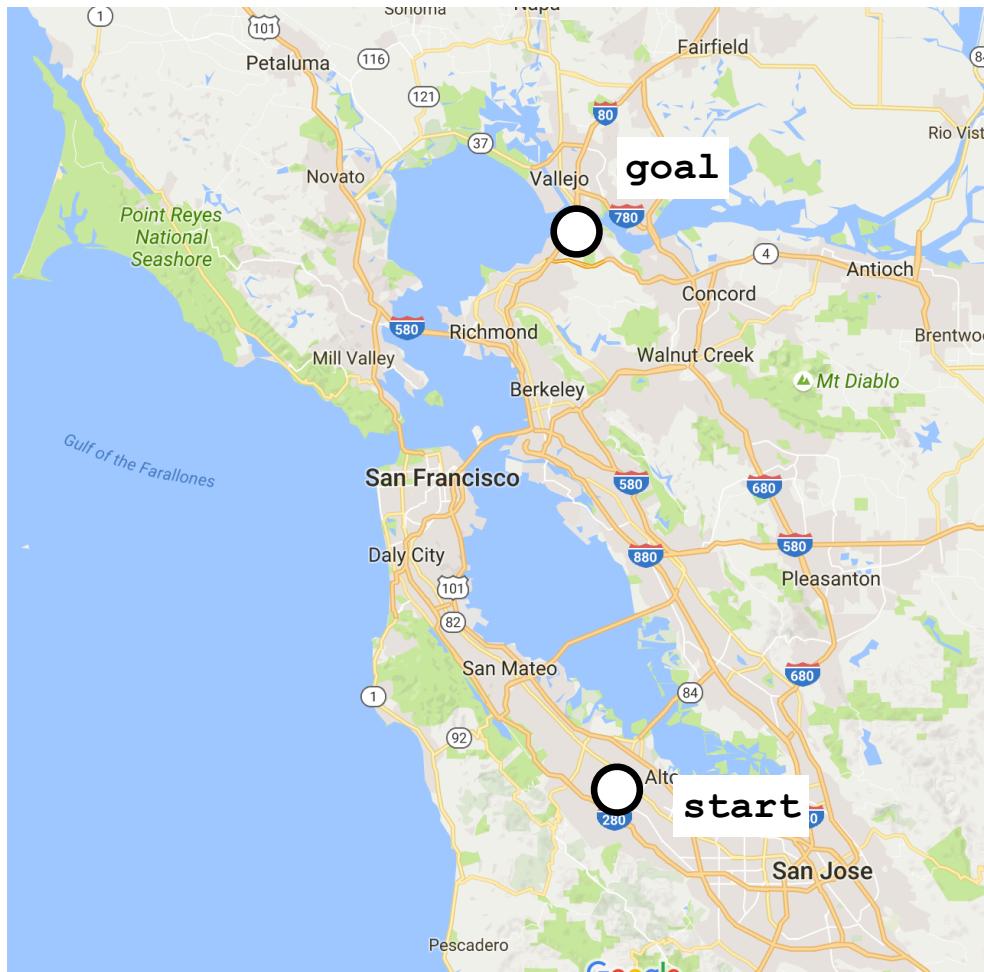
$$n^2$$

6×10^{15}

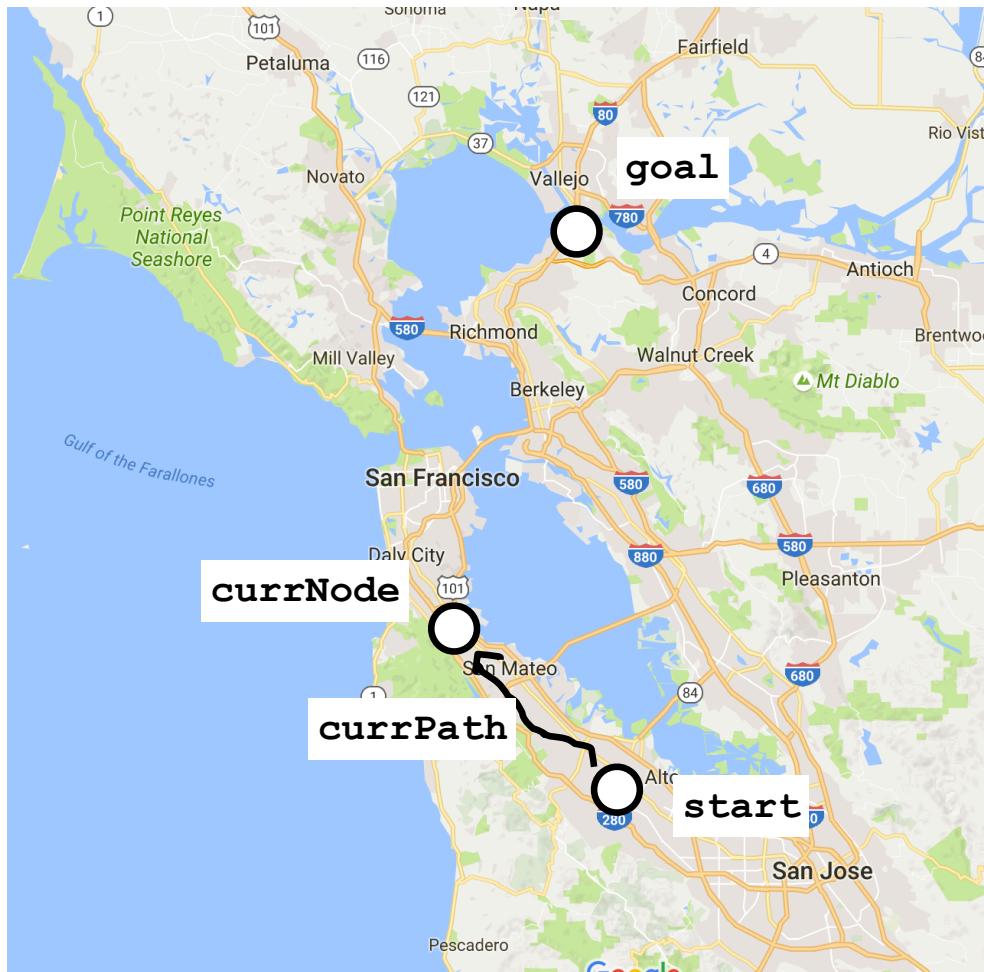
1 petasecond = 31.7 million years

Can you think of a heuristic?

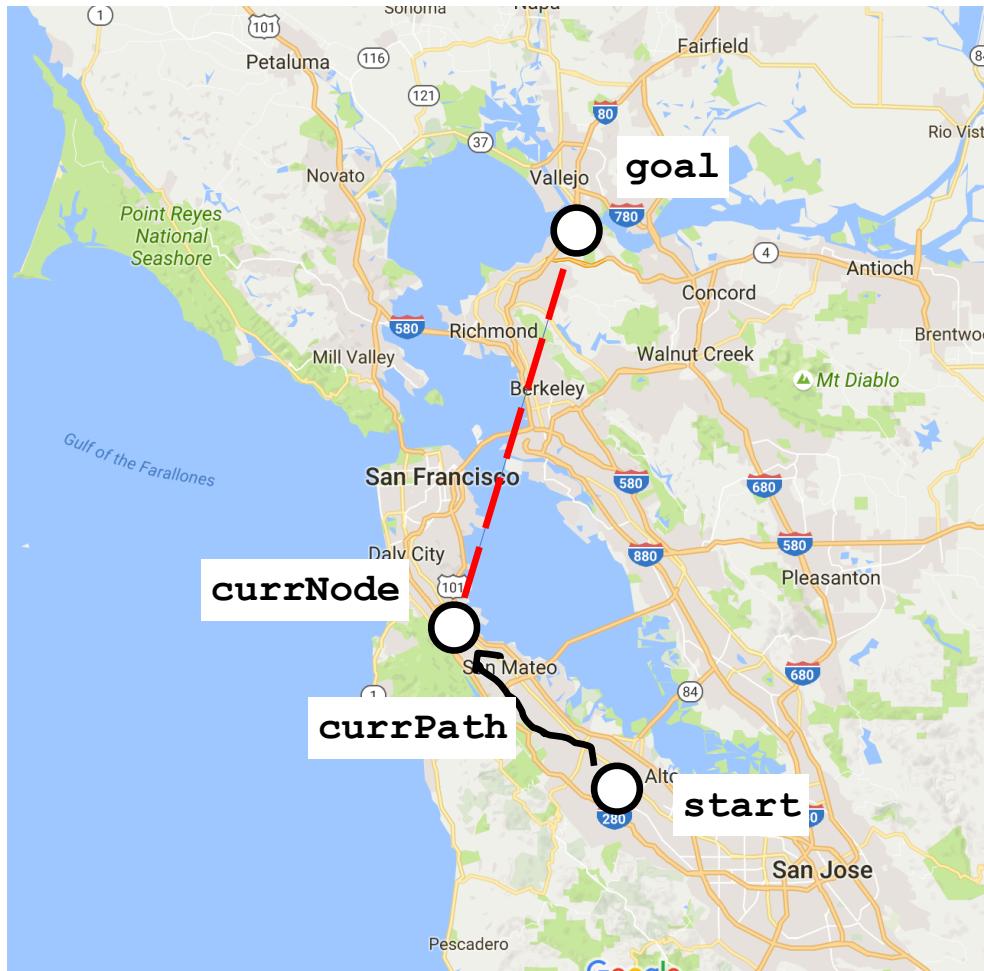
Road Map Heuristic



Road Map Heuristic

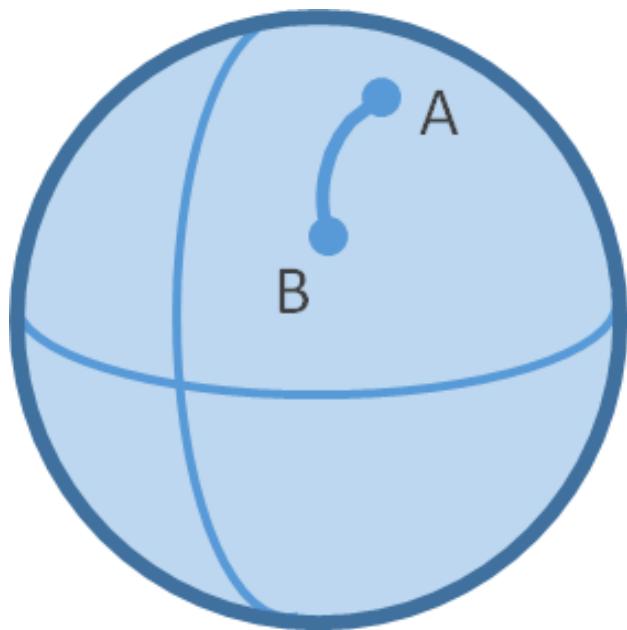


We must underestimate this time



Direct Highway

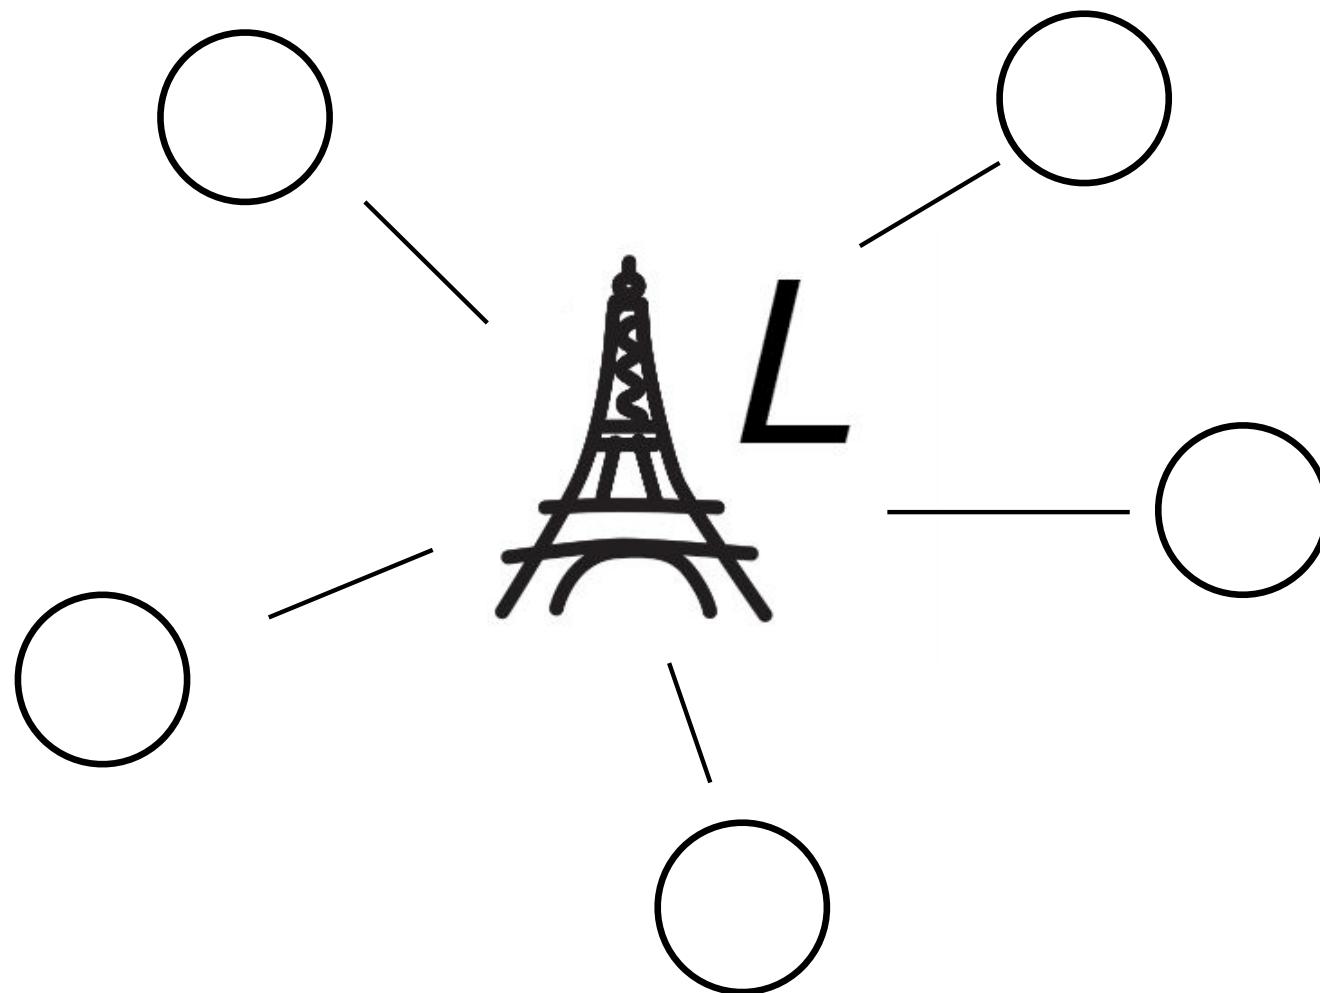
$$\text{Heuristic} = \frac{\text{Distance on surface of earth}}{\text{Speed on fastest highway}}$$



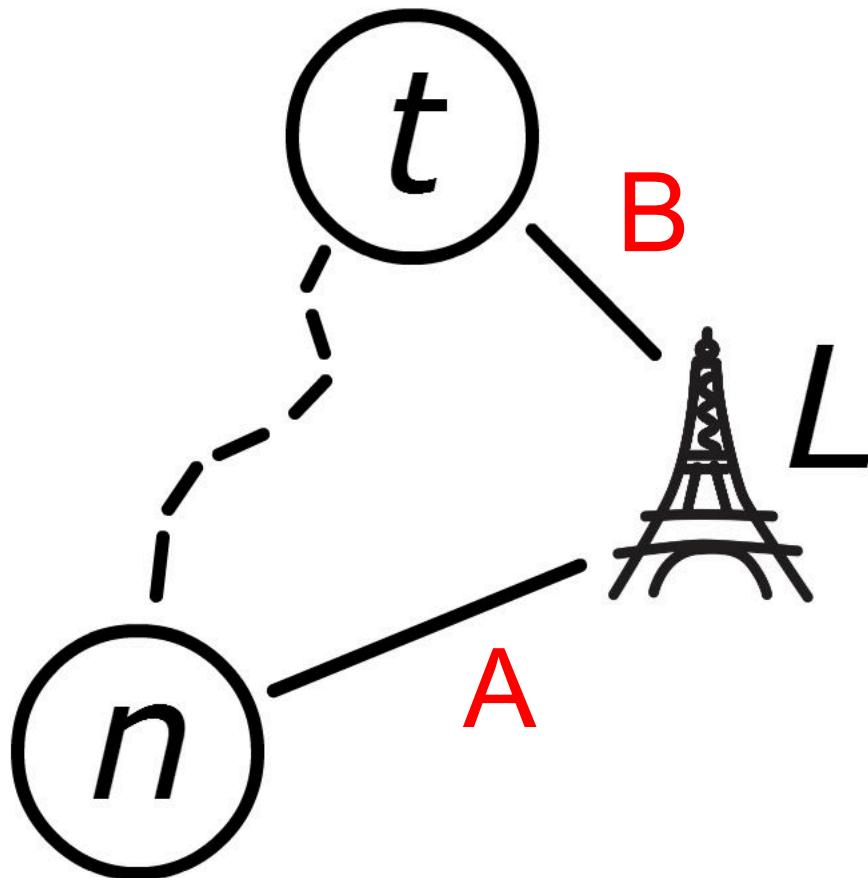
For Trailblazer:

Distance on surface of earth is `getCrowFlyDistance()`
Speed on fastest highway is `getMaxRoadSpeed()`

Distance to Landmarks



Landmark Heuristic



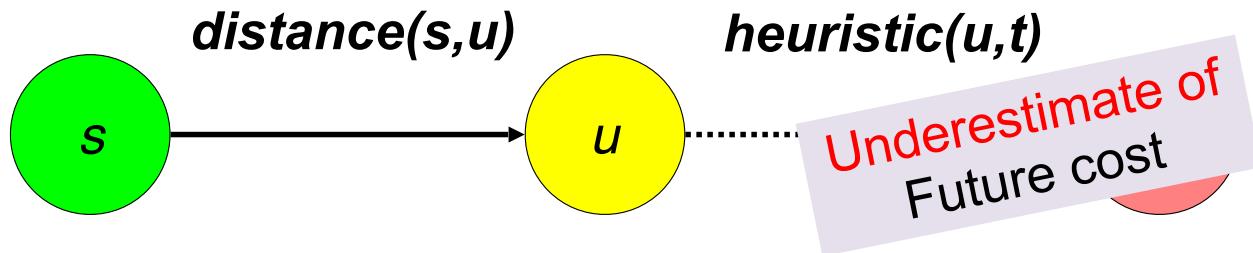
Distance < $\text{abs}(A - B)$

Best of All Heuristics

$$h = \max(h_1, h_2, \dots, h_n)$$

More Detail on A*: Choice of Heuristic

$$\text{priority}(u) = \text{distance}(s, u) + \text{heuristic}(u, t)$$



We want to underestimate the cost of our heuristic, by why?

Let's look at the bounds of our choices:

$$\text{heuristic}(u, t) = 0$$

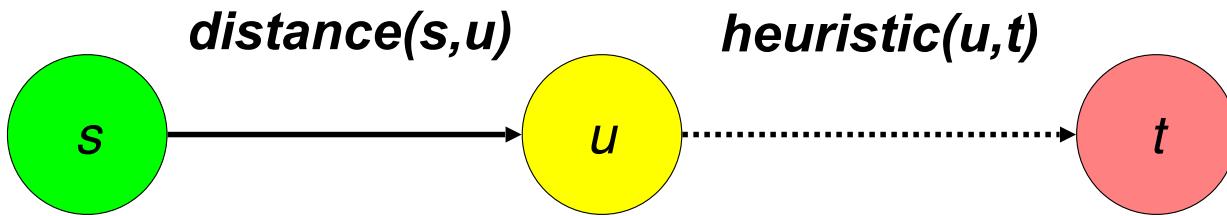
$$\text{heuristic}(u, t) = \text{underestimate}$$

$$\text{heuristic}(u, t) = \text{perfect distance}$$

$$\text{heuristic}(u, t) = \text{overestimate}$$

More Detail on A*: Choice of Heuristic

$$\text{priority}(u) = \text{distance}(s, u) + \text{heuristic}(u, t)$$



We want to underestimate the cost of our heuristic, by why?

Let's look at the bounds of our choices:

$\text{heuristic}(u,t) = 0$

$\text{heuristic}(u,t) = \text{underestimate}$

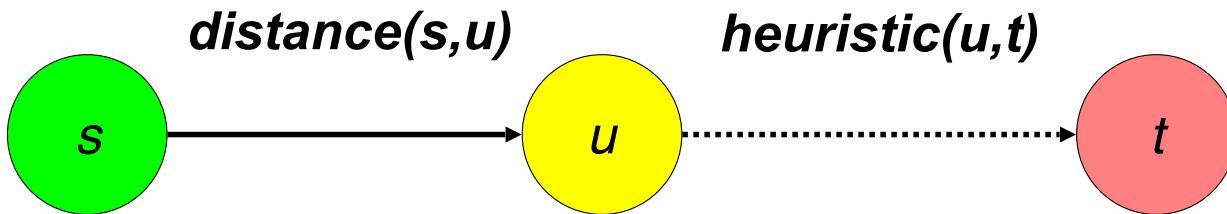
$\text{heuristic}(u,t) = \text{perfect distance}$

$\text{heuristic}(u,t) = \text{overestimate}$

Same as Dijkstra

More Detail on A*: Choice of Heuristic

$$\text{priority}(u) = \text{distance}(s, u) + \text{heuristic}(u, t)$$



We want to underestimate the cost of our heuristic, by why?

Let's look at the bounds of our choices:

$\text{heuristic}(u,t) = 0$

$\text{heuristic}(u,t) = \text{underestimate}$

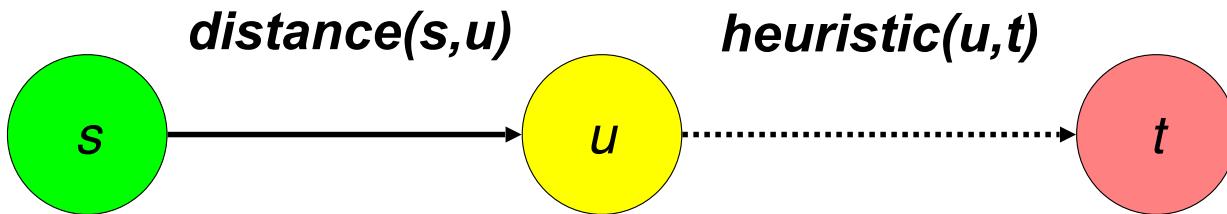
$\text{heuristic}(u,t) = \text{perfect distance}$

$\text{heuristic}(u,t) = \text{overestimate}$

Will be the same or faster than Dijkstra, and will find the shortest path (this is the only "admissible" heuristic for A*).

More Detail on A*: Choice of Heuristic

$$\text{priority}(u) = \text{distance}(s, u) + \text{heuristic}(u, t)$$



We want to underestimate the cost of our heuristic, by why?

Let's look at the bounds of our choices:

$\text{heuristic}(u,t) = 0$

$\text{heuristic}(u,t) = \text{underestimate}$

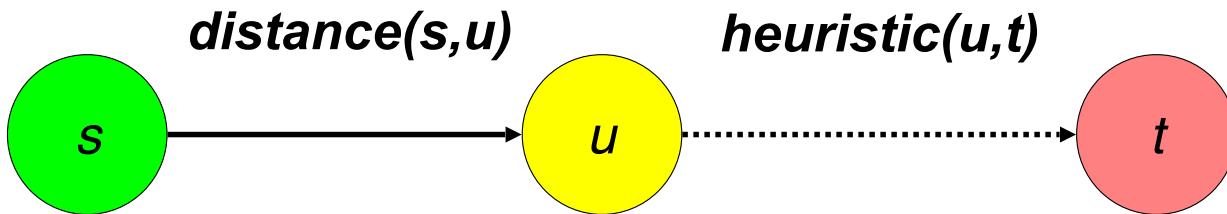
$\text{heuristic}(u,t) = \text{perfect distance}$

$\text{heuristic}(u,t) = \text{overestimate}$

Will only follow the best path, and will find the best path fastest (but requires perfect knowledge)

More Detail on A*: Choice of Heuristic

$$\text{priority}(u) = \text{distance}(s, u) + \text{heuristic}(u, t)$$



We want to underestimate the cost of our heuristic, by why?

Let's look at the bounds of our choices:

$\text{heuristic}(u,t) = 0$

$\text{heuristic}(u,t) = \text{underestimate}$

$\text{heuristic}(u,t) = \text{perfect distance}$

$\text{heuristic}(u,t) = \text{overestimate}$

**Won't necessarily find
shortest path (but might run
even faster)**

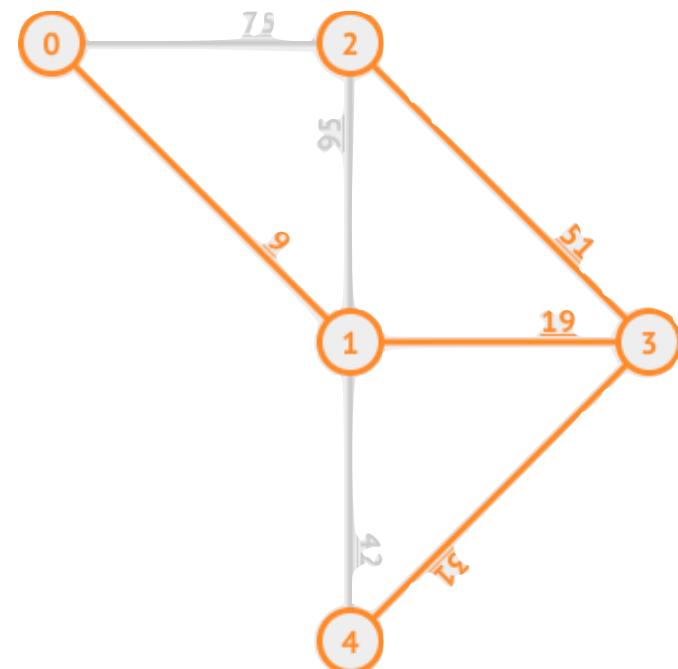
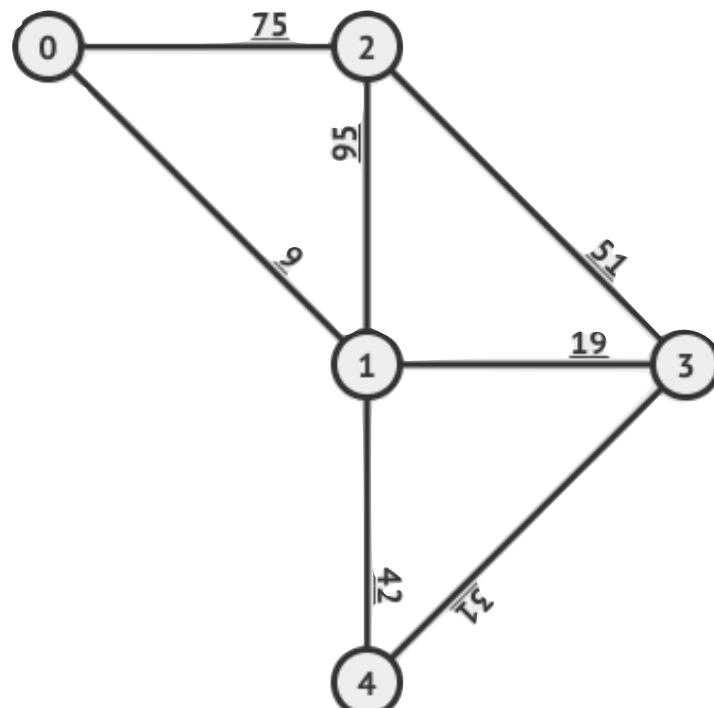
Admissible Heuristic

Definition: An admissible heuristic always **underestimates** the true cost.

Could you precompute this for all your vertices? Yes, but it would not be feasible.

Spanning Trees and Minimum Spanning Trees

Definition: A **Spanning Tree (ST)** of a connected undirected weighted graph **G** is a subgraph of **G** that is a **tree** and **connects (spans) all vertices of G**. A graph **G** can have multiple STs. A **Minimum Spanning Tree (MST)** of **G** is a ST of **G** that has the **smallest total weight** among the various STs. A graph **G** can have multiple MSTs but the MST weight is unique.



Minimum Spanning Tree

Kruskal's Algorithm to find a Minimum Spanning Tree

- **Kruskal's algorithm:** Finds a MST in a given graph.

function **kruskal**(graph):

 Remove all edges from the graph.

 Place all edges into a **priority queue** based on their weight (cost).

 While the priority queue is not empty:

 Dequeue an edge e from the priority queue.

 If e 's endpoints aren't already connected to one another,
 add that edge into the graph.

 Otherwise, skip the edge.

Kruskal Example

- In what order would Kruskal's algorithm visit the edges in the graph below? What MST would it produce?

```
function kruskal(graph):
```

Remove all edges from the graph.

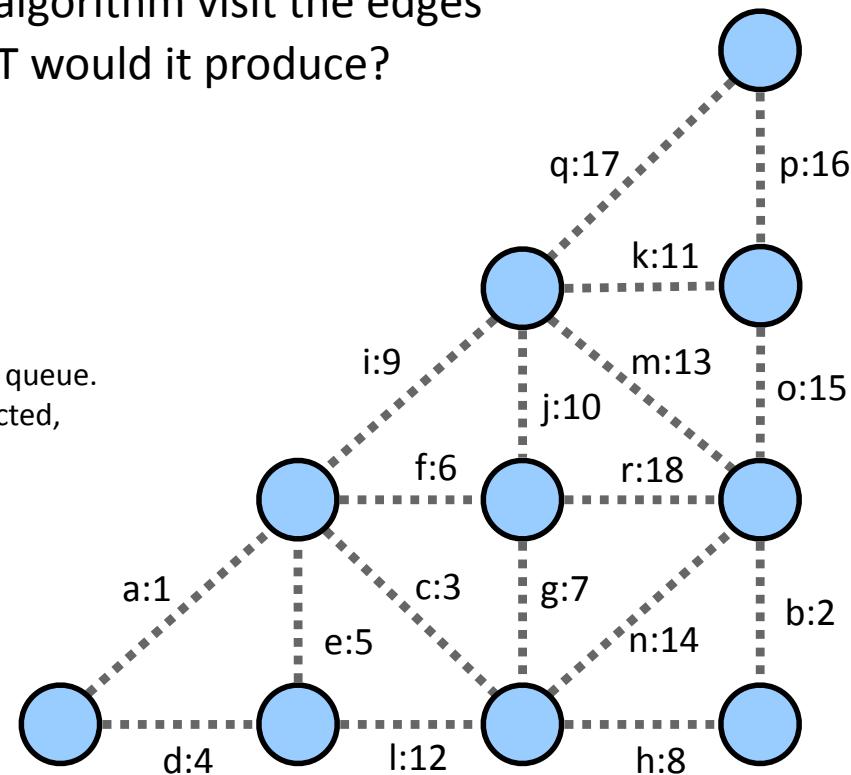
Place all edges into a priority queue based on their weight (cost).

While the priority queue is not empty:

Dequeue an edge e from the priority queue.

If e 's endpoints aren't already connected,
 add that edge into the graph.

Otherwise, skip the edge.

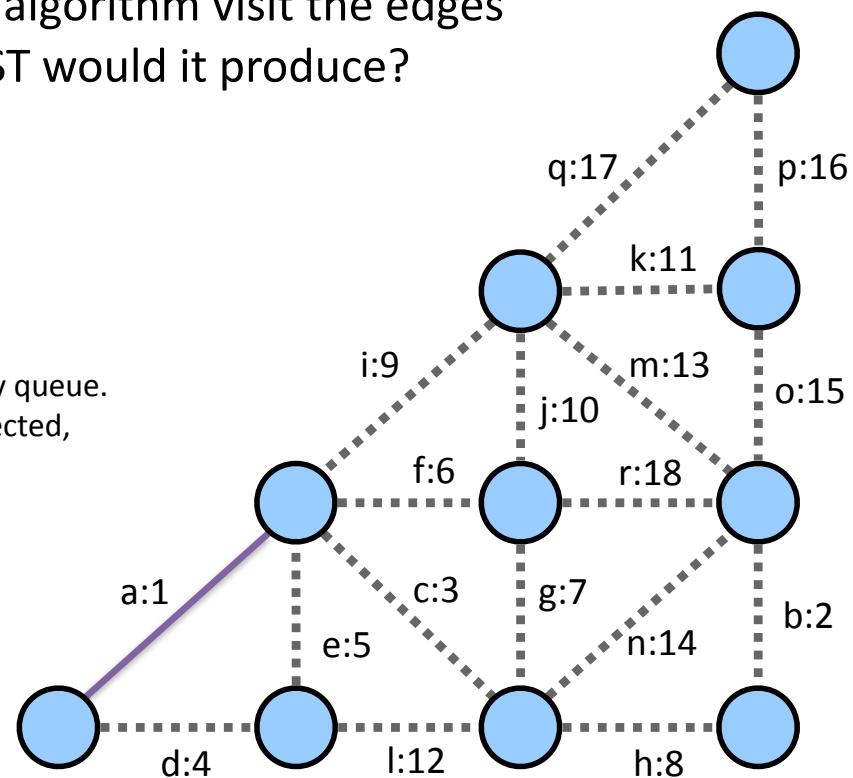


```
pq = {a:1, b:2, c:3, d:4, e:5, f:6, g:7, h:8, i:9, j:10, k:11, l:12, m:13, n:14, o:15, p:16, q:17, r:18}
```

Kruskal Example

- In what order would Kruskal's algorithm visit the edges in the graph below? What MST would it produce?

```
function kruskal(graph):
    Remove all edges from the graph.
    Place all edges into a priority queue
        based on their weight (cost).
    While the priority queue is not empty:
        Dequeue an edge e from the priority queue.
        If e's endpoints aren't already connected,
            add that edge into the graph.
        Otherwise, skip the edge.
```

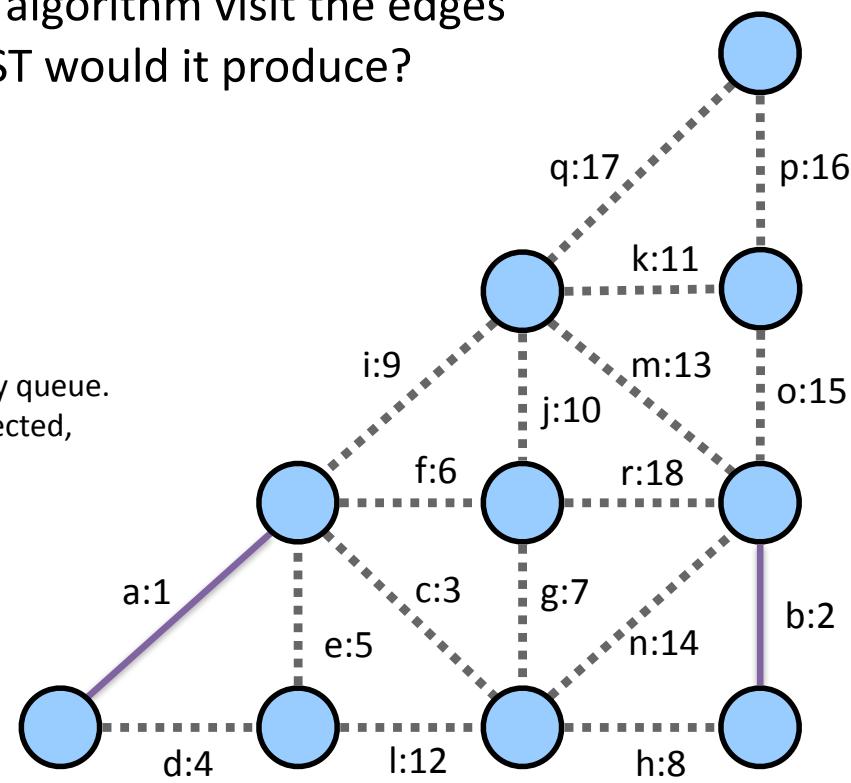


$\text{pq} = \{\mathbf{a:1}, b:2, c:3, d:4, e:5, f:6, g:7, h:8, i:9, j:10, k:11, l:12, m:13, n:14, o:15, p:16, q:17, r:18\}$

Kruskal Example

- In what order would Kruskal's algorithm visit the edges in the graph below? What MST would it produce?

```
function kruskal(graph):
    Remove all edges from the graph.
    Place all edges into a priority queue
        based on their weight (cost).
    While the priority queue is not empty:
        Dequeue an edge e from the priority queue.
        If e's endpoints aren't already connected,
            add that edge into the graph.
        Otherwise, skip the edge.
```

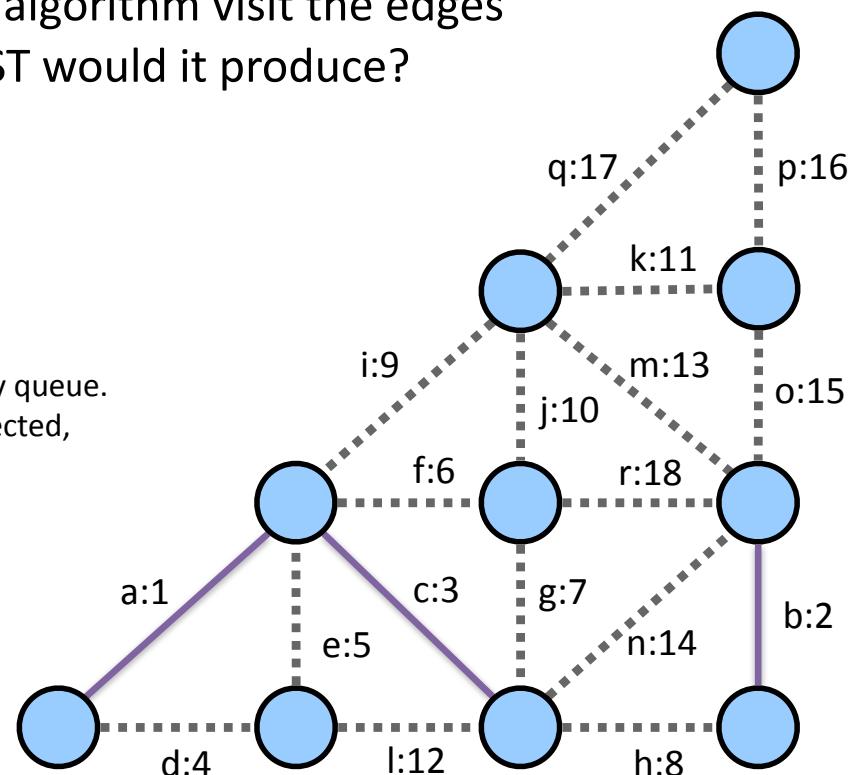


$\text{pq} = \{\mathbf{b:2}, c:3, d:4, e:5, f:6, g:7, h:8, i:9, j:10, k:11, l:12, m:13, n:14, o:15, p:16, q:17, r:18\}$

Kruskal Example

- In what order would Kruskal's algorithm visit the edges in the graph below? What MST would it produce?

```
function kruskal(graph):
    Remove all edges from the graph.
    Place all edges into a priority queue
        based on their weight (cost).
    While the priority queue is not empty:
        Dequeue an edge e from the priority queue.
        If e's endpoints aren't already connected,
            add that edge into the graph.
        Otherwise, skip the edge.
```

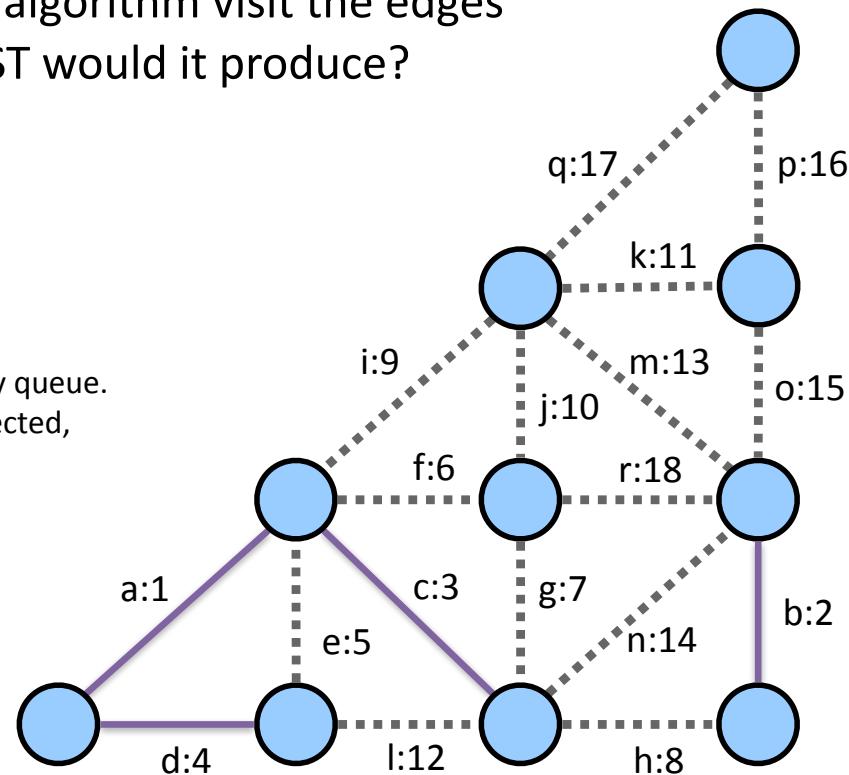


$\text{pq} = \{\mathbf{C:3}, d:4, e:5, f:6, g:7, h:8, i:9, j:10, k:11, l:12, m:13, n:14, o:15, p:16, q:17, r:18\}$

Kruskal Example

- In what order would Kruskal's algorithm visit the edges in the graph below? What MST would it produce?

```
function kruskal(graph):
    Remove all edges from the graph.
    Place all edges into a priority queue
        based on their weight (cost).
    While the priority queue is not empty:
        Dequeue an edge e from the priority queue.
        If e's endpoints aren't already connected,
            add that edge into the graph.
        Otherwise, skip the edge.
```

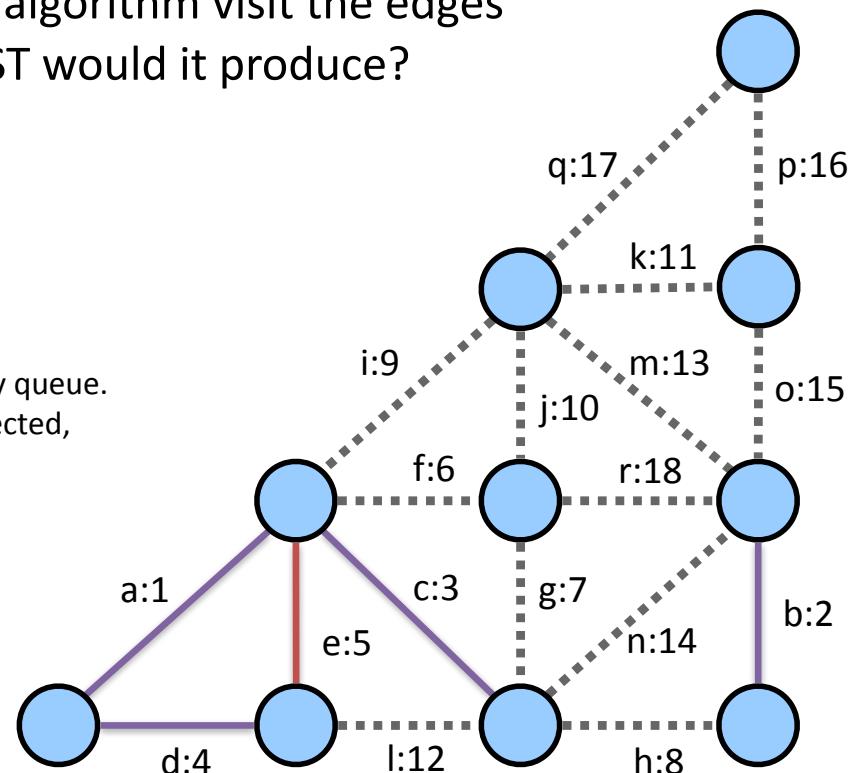


pq = {d:4, e:5, f:6, g:7, h:8, i:9, j:10, k:11, l:12, m:13, n:14, o:15, p:16, q:17, r:18}

Kruskal Example

- In what order would Kruskal's algorithm visit the edges in the graph below? What MST would it produce?

```
function kruskal(graph):
    Remove all edges from the graph.
    Place all edges into a priority queue
        based on their weight (cost).
    While the priority queue is not empty:
        Dequeue an edge e from the priority queue.
        If e's endpoints aren't already connected,
            add that edge into the graph.
        Otherwise, skip the edge.
```

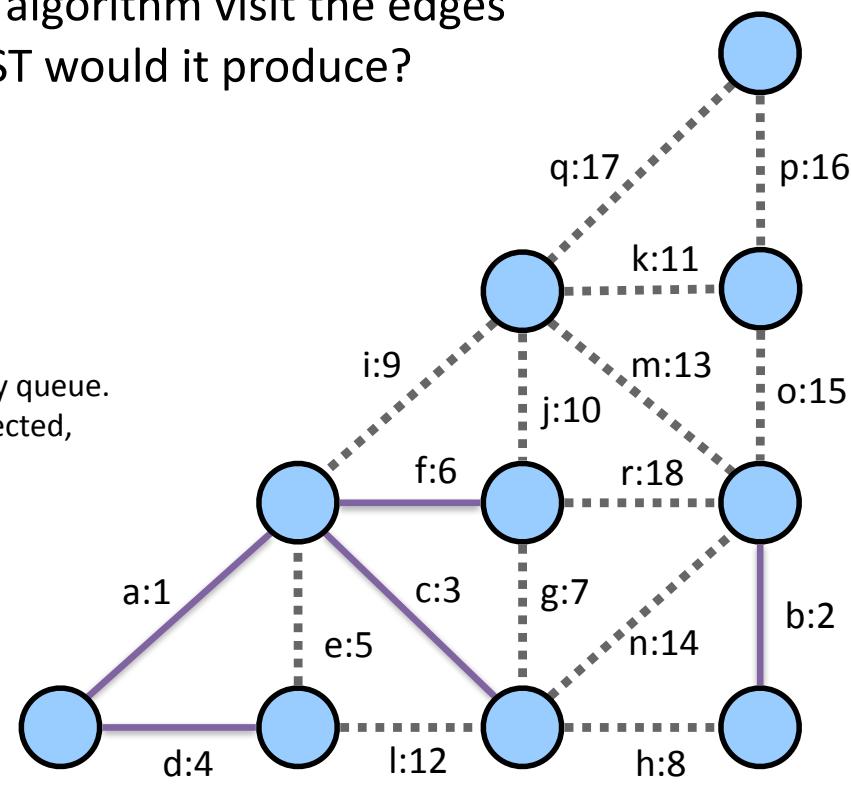


$pq = \{e:5, f:6, g:7, h:8, i:9, j:10, k:11, l:12, m:13, n:14, o:15, p:16, q:17, r:18\}$

Kruskal Example

- In what order would Kruskal's algorithm visit the edges in the graph below? What MST would it produce?

```
function kruskal(graph):
    Remove all edges from the graph.
    Place all edges into a priority queue
        based on their weight (cost).
    While the priority queue is not empty:
        Dequeue an edge e from the priority queue.
        If e's endpoints aren't already connected,
            add that edge into the graph.
        Otherwise, skip the edge.
```

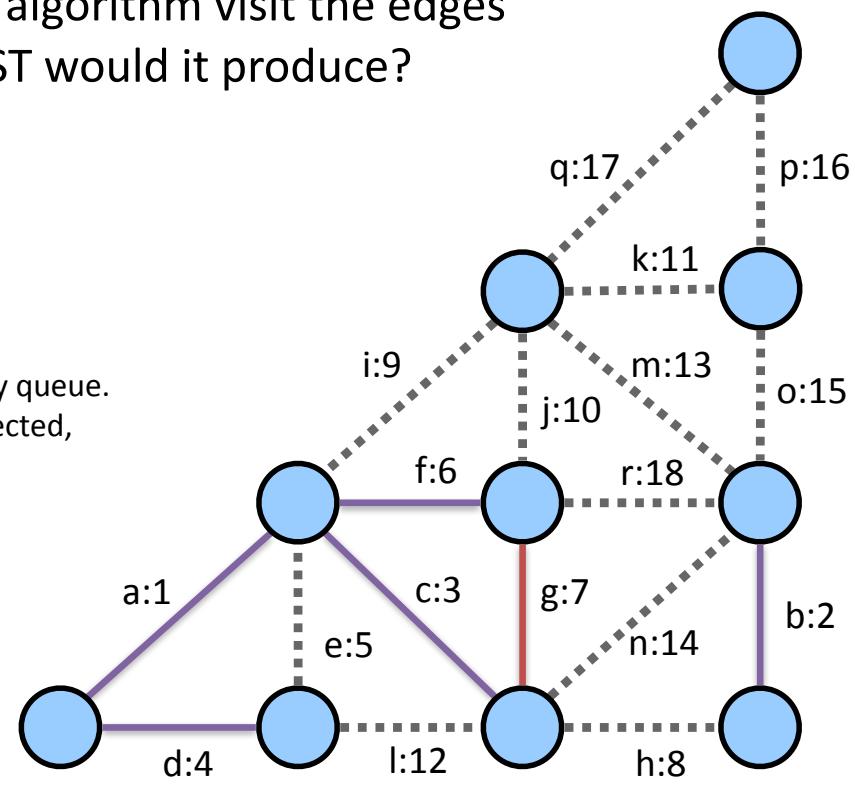


$\text{pq} = \{\mathbf{f:6}, g:7, h:8, i:9, j:10, k:11, l:12, m:13, n:14, o:15, p:16, q:17, r:18\}$

Kruskal Example

- In what order would Kruskal's algorithm visit the edges in the graph below? What MST would it produce?

```
function kruskal(graph):
    Remove all edges from the graph.
    Place all edges into a priority queue
        based on their weight (cost).
    While the priority queue is not empty:
        Dequeue an edge e from the priority queue.
        If e's endpoints aren't already connected,
            add that edge into the graph.
        Otherwise, skip the edge.
```

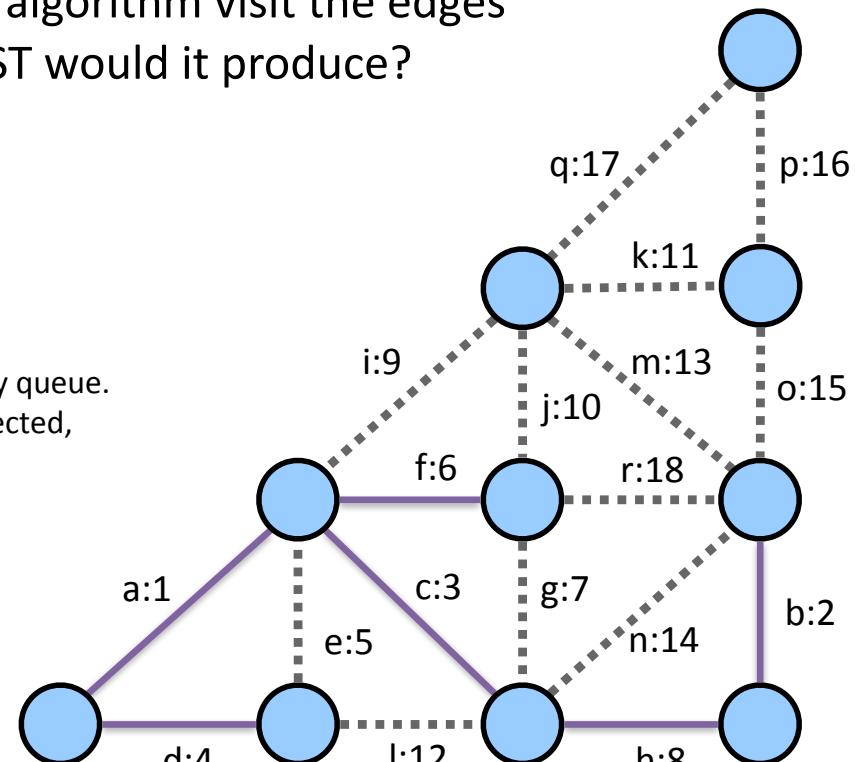


$\text{pq} = \{\mathbf{g:7}, h:8, i:9, j:10, k:11, l:12, m:13, n:14, o:15, p:16, q:17, r:18\}$

Kruskal Example

- In what order would Kruskal's algorithm visit the edges in the graph below? What MST would it produce?

```
function kruskal(graph):
    Remove all edges from the graph.
    Place all edges into a priority queue
        based on their weight (cost).
    While the priority queue is not empty:
        Dequeue an edge e from the priority queue.
        If e's endpoints aren't already connected,
            add that edge into the graph.
        Otherwise, skip the edge.
```



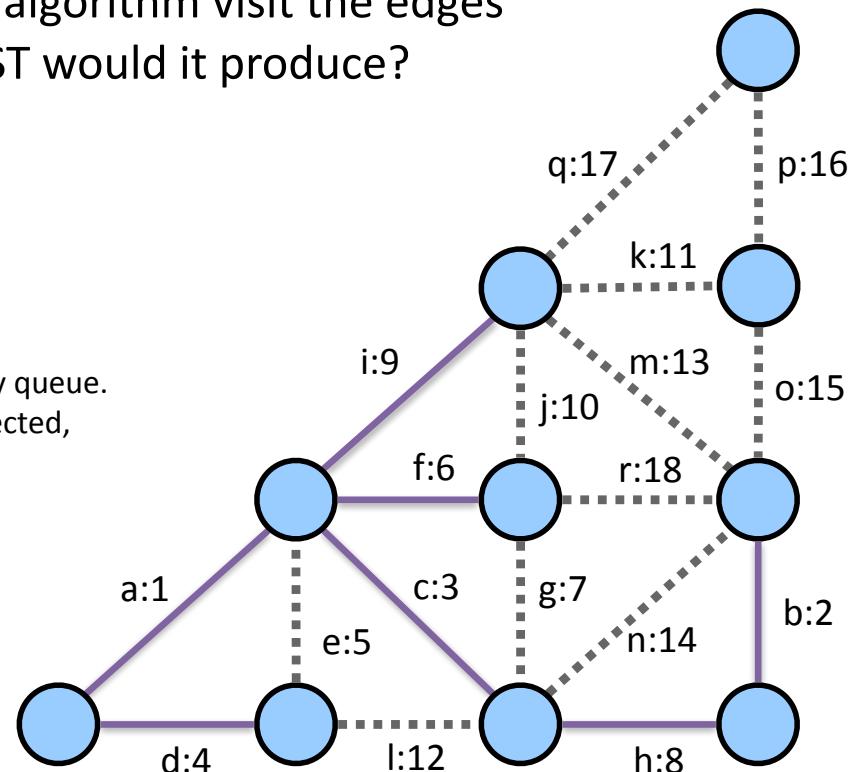
$\text{pq} = \{\text{h:8}, \text{i:9}, \text{j:10}, \text{k:11}, \text{l:12}, \text{m:13}, \text{n:14}, \text{o:15}, \text{p:16}, \text{q:17}, \text{r:18}\}$

Kruskal Example

- In what order would Kruskal's algorithm visit the edges in the graph below? What MST would it produce?

```
function kruskal(graph):
    Remove all edges from the graph.
    Place all edges into a priority queue
        based on their weight (cost).
    While the priority queue is not empty:
        Dequeue an edge e from the priority queue.
        If e's endpoints aren't already connected,
            add that edge into the graph.
        Otherwise, skip the edge.
```

$pq = \{i:9, j:10, k:11, l:12, m:13, n:14, o:15, p:16, q:17, r:18\}$

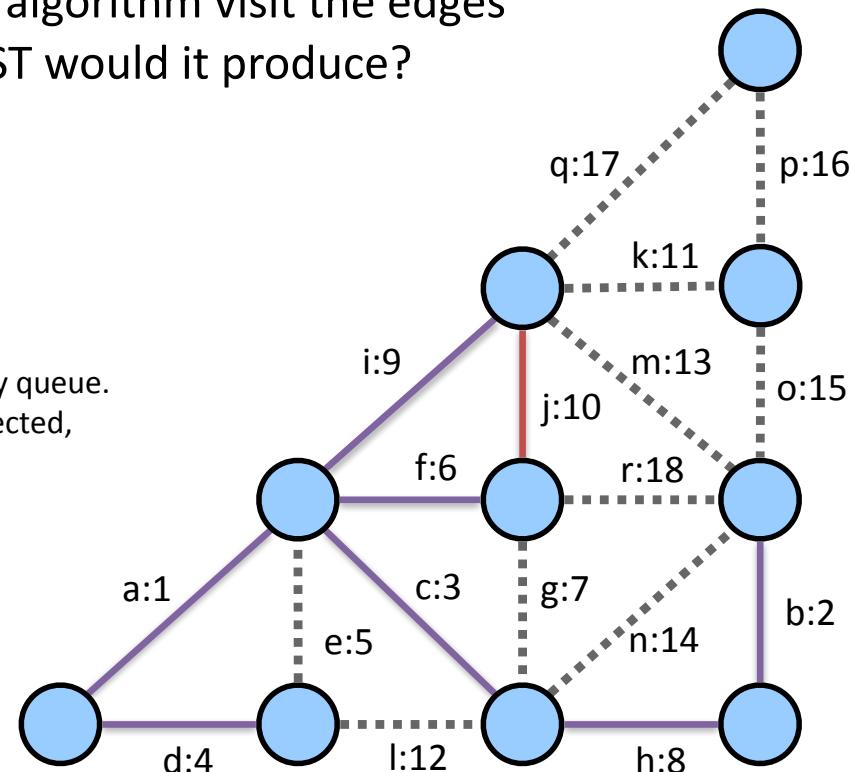


Kruskal Example

- In what order would Kruskal's algorithm visit the edges in the graph below? What MST would it produce?

```
function kruskal(graph):
    Remove all edges from the graph.
    Place all edges into a priority queue
        based on their weight (cost).
    While the priority queue is not empty:
        Dequeue an edge e from the priority queue.
        If e's endpoints aren't already connected,
            add that edge into the graph.
        Otherwise, skip the edge.
```

$pq = \{j:10, k:11, l:12, m:13, n:14, o:15, p:16, q:17, r:18\}$

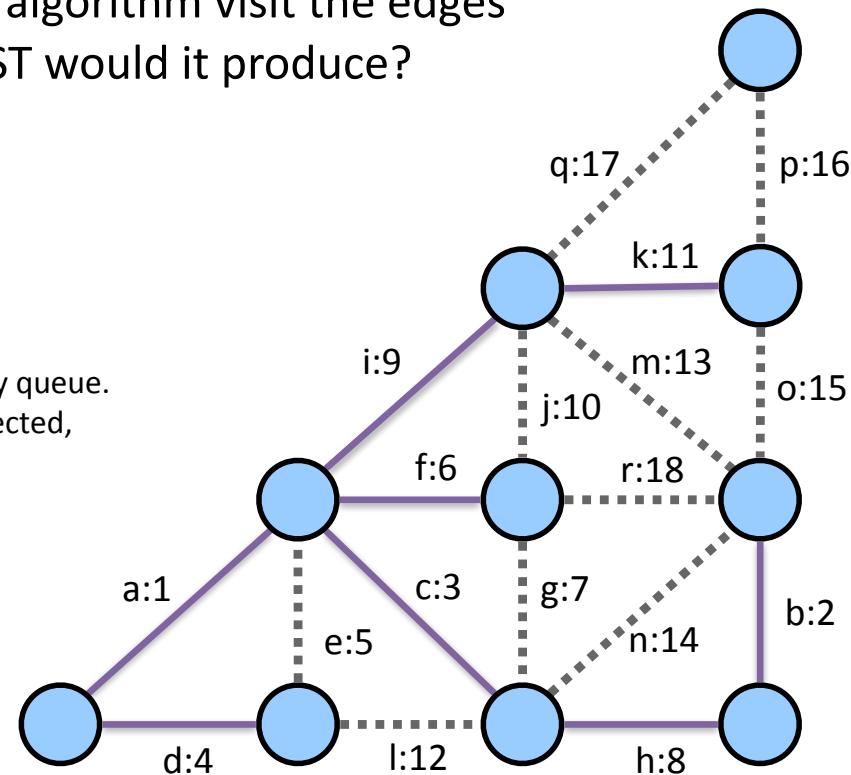


Kruskal Example

- In what order would Kruskal's algorithm visit the edges in the graph below? What MST would it produce?

```
function kruskal(graph):
    Remove all edges from the graph.
    Place all edges into a priority queue
        based on their weight (cost).
    While the priority queue is not empty:
        Dequeue an edge e from the priority queue.
        If e's endpoints aren't already connected,
            add that edge into the graph.
        Otherwise, skip the edge.
```

pq = {k:11, l:12, m:13, n:14, o:15, p:16, q:17, r:18}

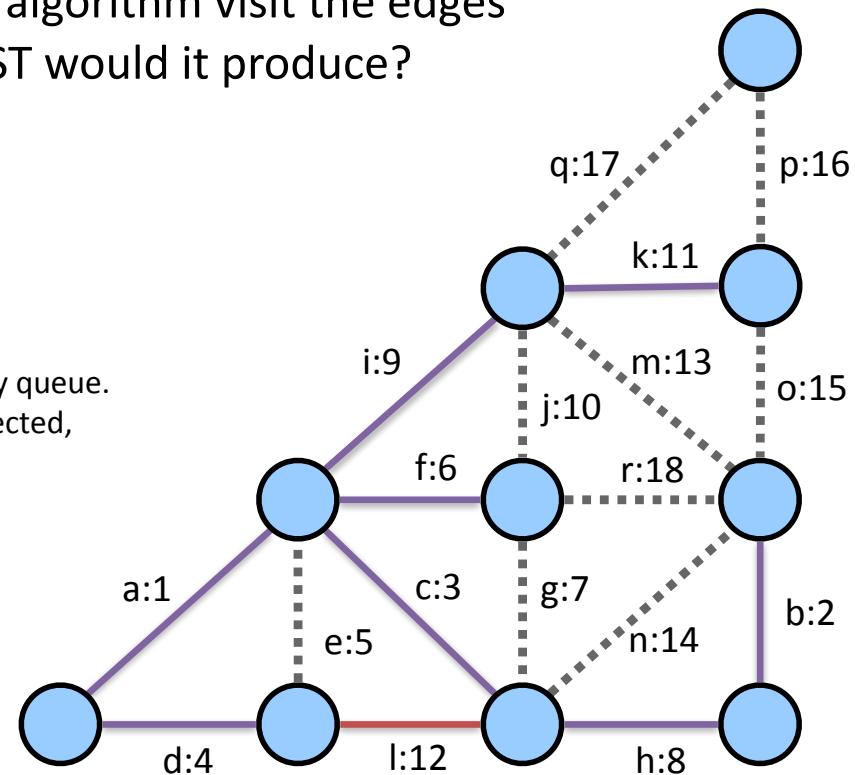


Kruskal Example

- In what order would Kruskal's algorithm visit the edges in the graph below? What MST would it produce?

```
function kruskal(graph):
    Remove all edges from the graph.
    Place all edges into a priority queue
        based on their weight (cost).
    While the priority queue is not empty:
        Dequeue an edge e from the priority queue.
        If e's endpoints aren't already connected,
            add that edge into the graph.
        Otherwise, skip the edge.
```

$pq = \{l:12, m:13, n:14, o:15, p:16, q:17, r:18\}$

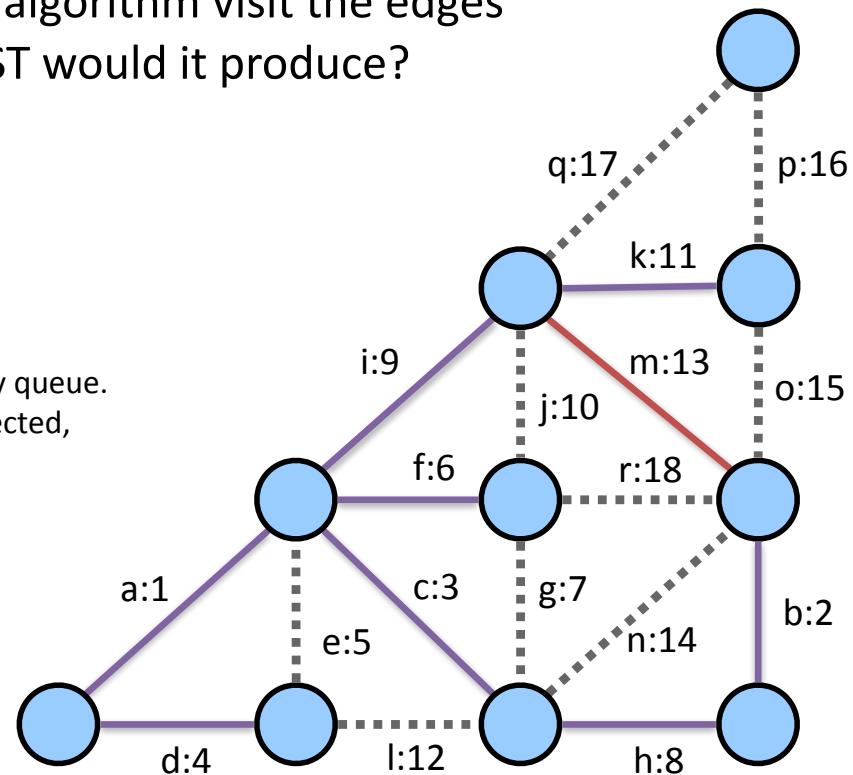


Kruskal Example

- In what order would Kruskal's algorithm visit the edges in the graph below? What MST would it produce?

```
function kruskal(graph):
    Remove all edges from the graph.
    Place all edges into a priority queue
        based on their weight (cost).
    While the priority queue is not empty:
        Dequeue an edge e from the priority queue.
        If e's endpoints aren't already connected,
            add that edge into the graph.
        Otherwise, skip the edge.
```

pq = {m:13, n:14, o:15, p:16, q:17, r:18}

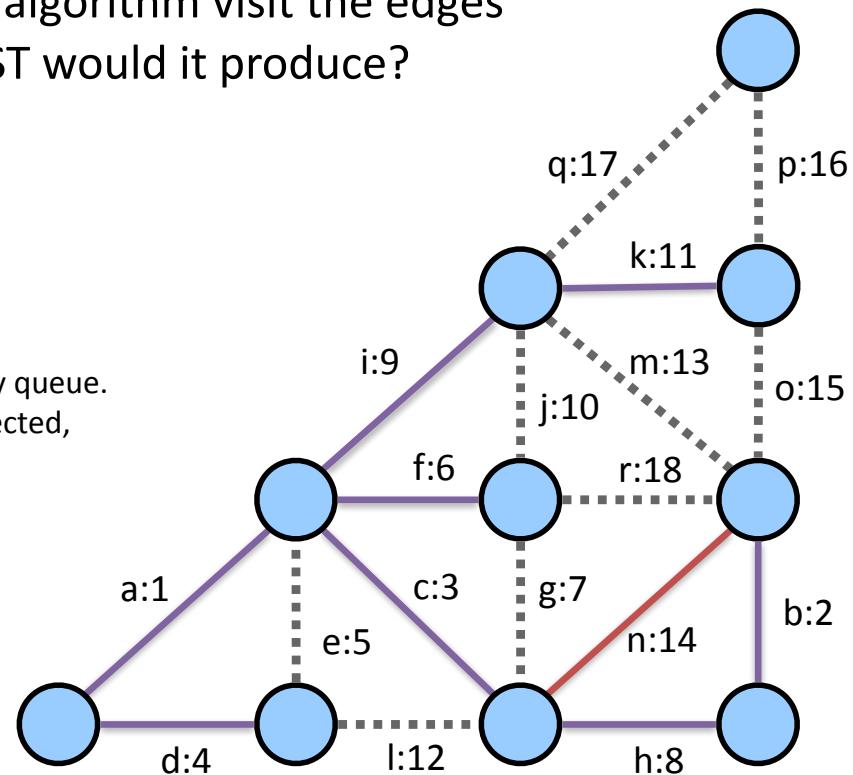


Kruskal Example

- In what order would Kruskal's algorithm visit the edges in the graph below? What MST would it produce?

```
function kruskal(graph):
    Remove all edges from the graph.
    Place all edges into a priority queue
        based on their weight (cost).
    While the priority queue is not empty:
        Dequeue an edge e from the priority queue.
        If e's endpoints aren't already connected,
            add that edge into the graph.
        Otherwise, skip the edge.
```

pq = {n:14, o:15, p:16, q:17, r:18}

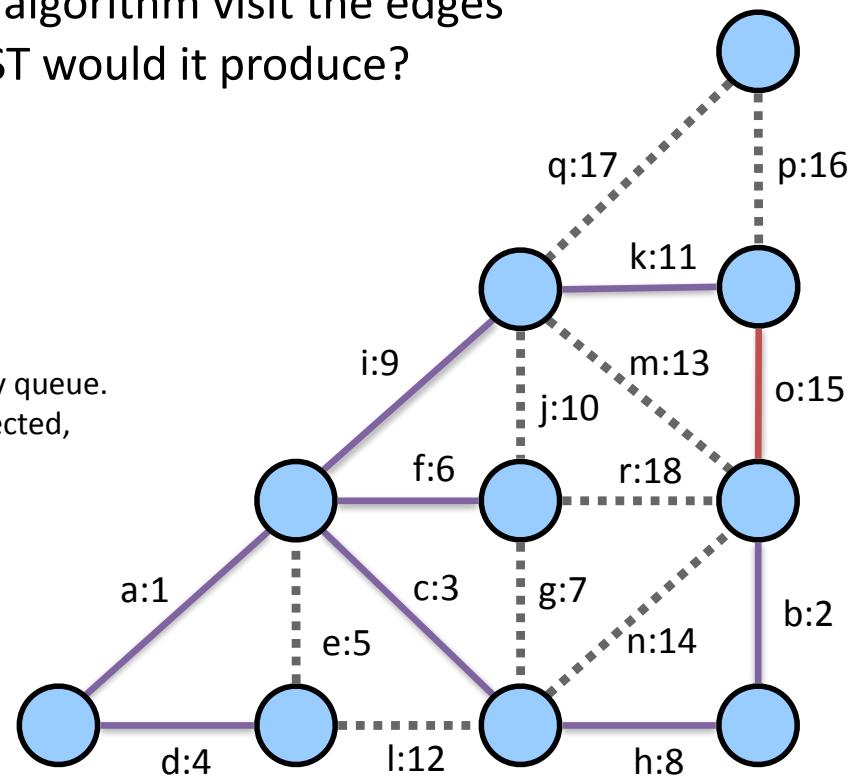


Kruskal Example

- In what order would Kruskal's algorithm visit the edges in the graph below? What MST would it produce?

```
function kruskal(graph):
    Remove all edges from the graph.
    Place all edges into a priority queue
        based on their weight (cost).
    While the priority queue is not empty:
        Dequeue an edge e from the priority queue.
        If e's endpoints aren't already connected,
            add that edge into the graph.
        Otherwise, skip the edge.
```

$pq = \{O:15, p:16, q:17, r:18\}$

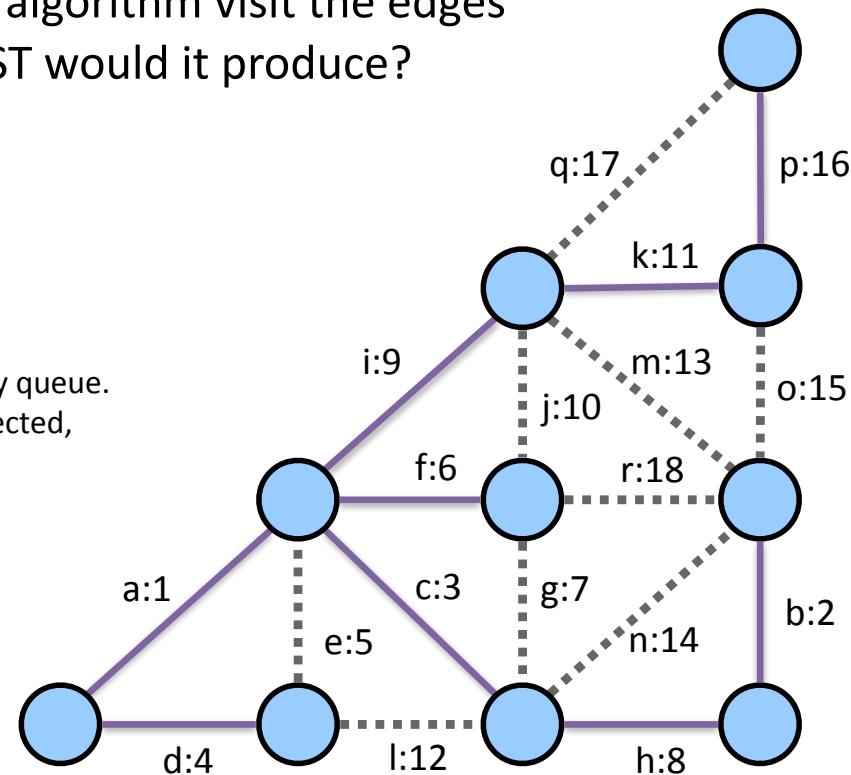


Kruskal Example

- In what order would Kruskal's algorithm visit the edges in the graph below? What MST would it produce?

```
function kruskal(graph):
    Remove all edges from the graph.
    Place all edges into a priority queue
        based on their weight (cost).
    While the priority queue is not empty:
        Dequeue an edge e from the priority queue.
        If e's endpoints aren't already connected,
            add that edge into the graph.
        Otherwise, skip the edge.
```

pq = {p:16, q:17, r:18}

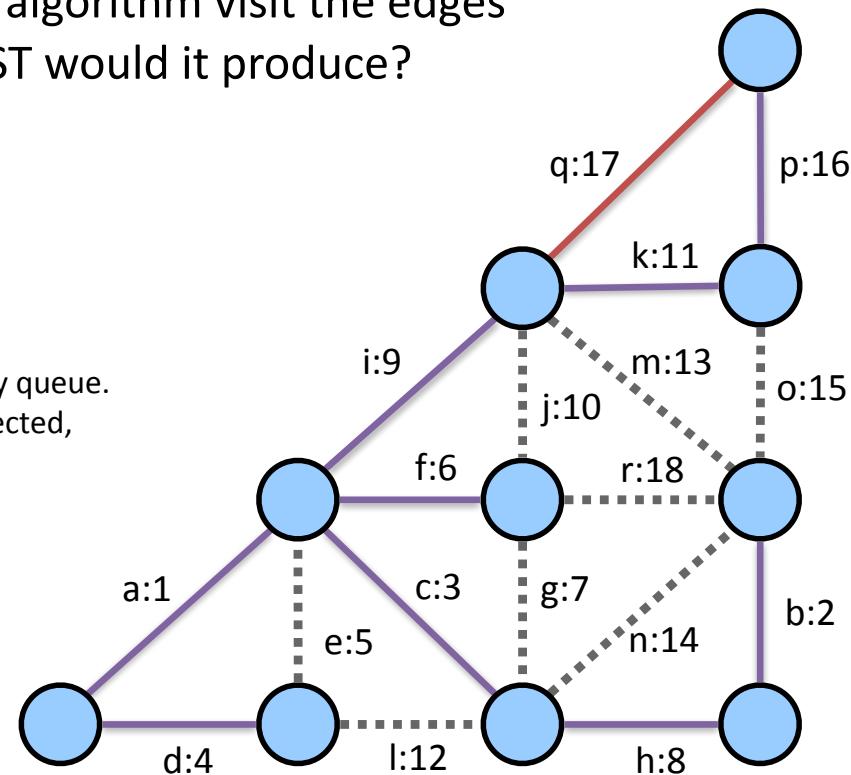


Kruskal Example

- In what order would Kruskal's algorithm visit the edges in the graph below? What MST would it produce?

```
function kruskal(graph):
    Remove all edges from the graph.
    Place all edges into a priority queue
        based on their weight (cost).
    While the priority queue is not empty:
        Dequeue an edge e from the priority queue.
        If e's endpoints aren't already connected,
            add that edge into the graph.
        Otherwise, skip the edge.
```

$pq = \{q:17, r:18\}$

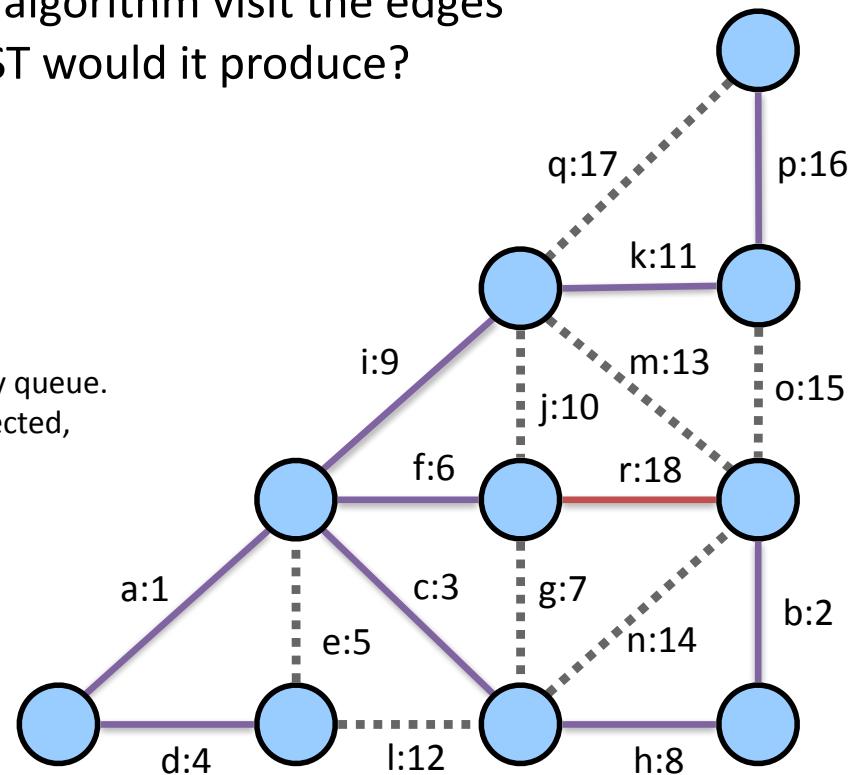


Kruskal Example

- In what order would Kruskal's algorithm visit the edges in the graph below? What MST would it produce?

```
function kruskal(graph):
    Remove all edges from the graph.
    Place all edges into a priority queue
        based on their weight (cost).
    While the priority queue is not empty:
        Dequeue an edge e from the priority queue.
        If e's endpoints aren't already connected,
            add that edge into the graph.
        Otherwise, skip the edge.
```

$pq = \{r:18\}$



Kruskal Example

- In what order would Kruskal's algorithm visit the edges in the graph below? What MST would it produce?

```
function kruskal(graph):
```

Remove all edges from the graph.

Place all edges into a priority queue
 based on their weight (cost).

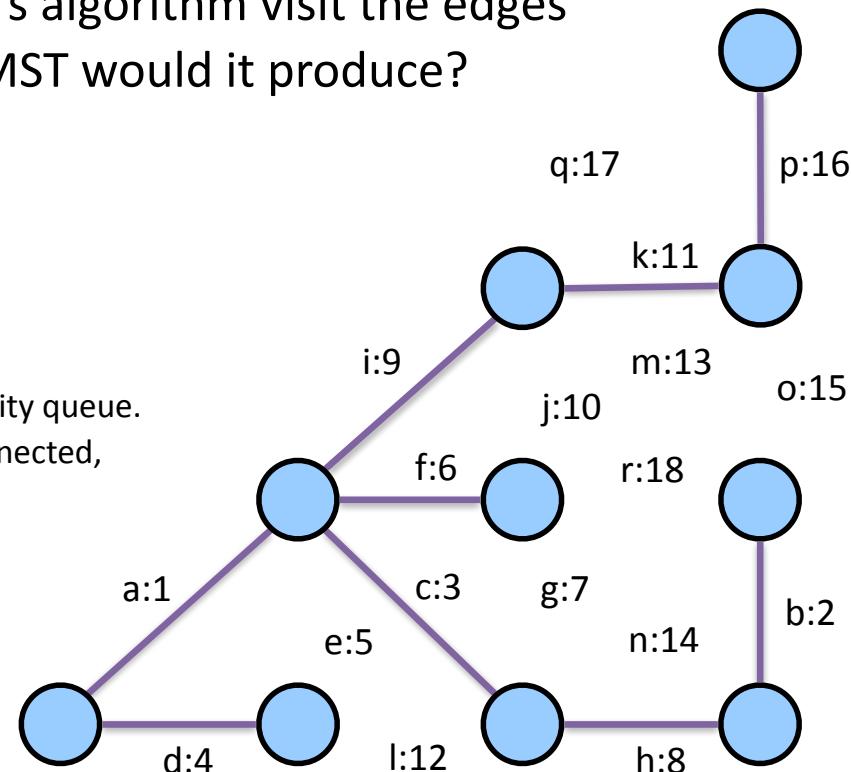
While the priority queue is not empty:

Dequeue an edge e from the priority queue.

If e 's endpoints aren't already connected,
 add that edge into the graph.

Otherwise, skip the edge.

$pq = \{\}$

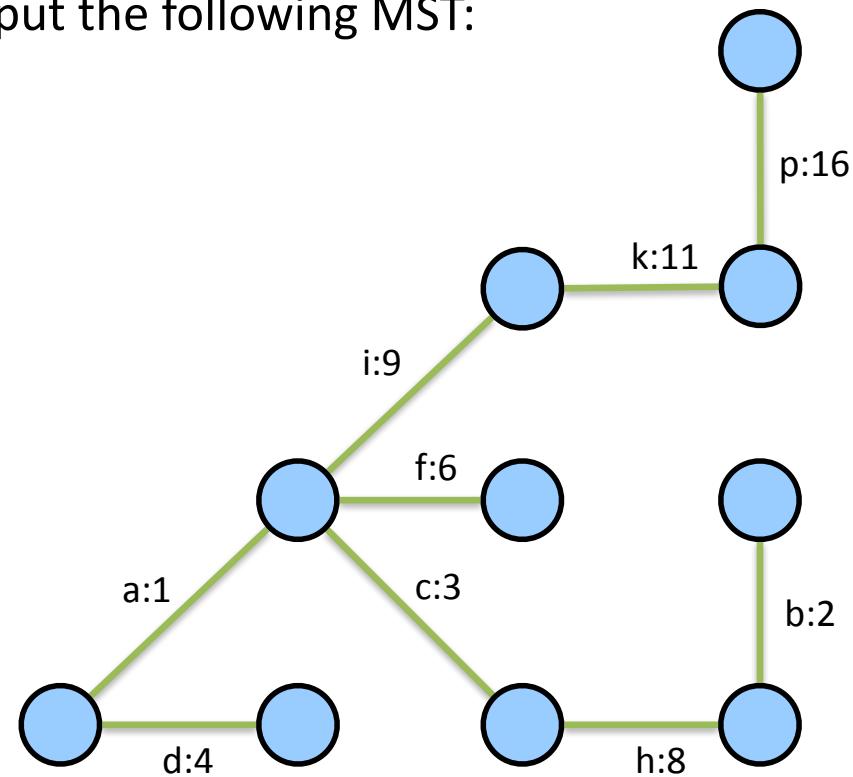


Kruskal Example

- Kruskal's algorithm would output the following MST:
 - {a, b, c, d, f, h, i, k, p}

- The MST's total cost is:

$$1+2+3+4+6+8+9+11+16 = 60$$



- What data structures should we use to implement this algorithm?

function **kruskal**(graph):

 Remove all edges from the graph.

 Place all edges into a **priority queue**
 based on their weight (cost).

 While the priority queue is not empty:

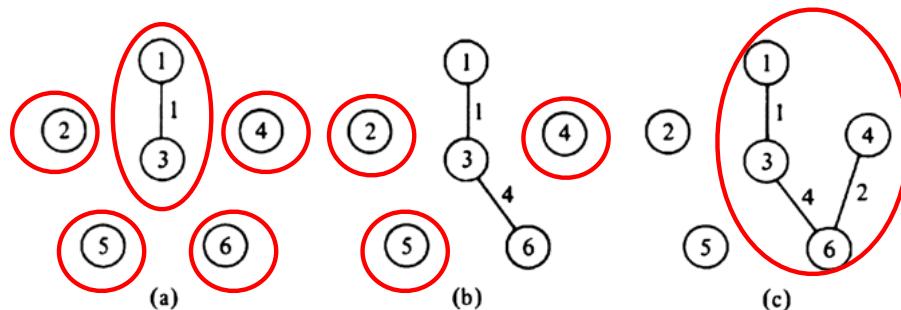
 Dequeue an edge e from the priority queue.

 If e 's endpoints aren't already connected,
 add that edge into the graph.

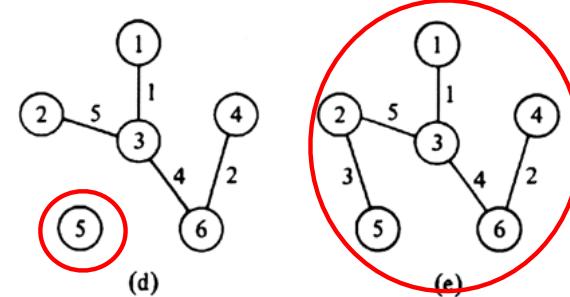
 Otherwise, skip the edge.

- Need some way to identify which vertexes are "connected" to which other ones
 - we call these "**clusters**" of vertices

- Also need an efficient way to figure out which cluster a given vertex is in.



- Also need to **merge clusters** when adding an edge.



References and Advanced Reading

- **References:**

- A* Heuristics: <http://theory.stanford.edu/~amitp/GameProgramming/Heuristics.html>
- Minimum Spanning Tree visualization: <https://visualgo.net/mst>
- Kruskal's Algorithm: https://en.wikipedia.org/wiki/Kruskal%27s_algorithm

- **Advanced Reading:**

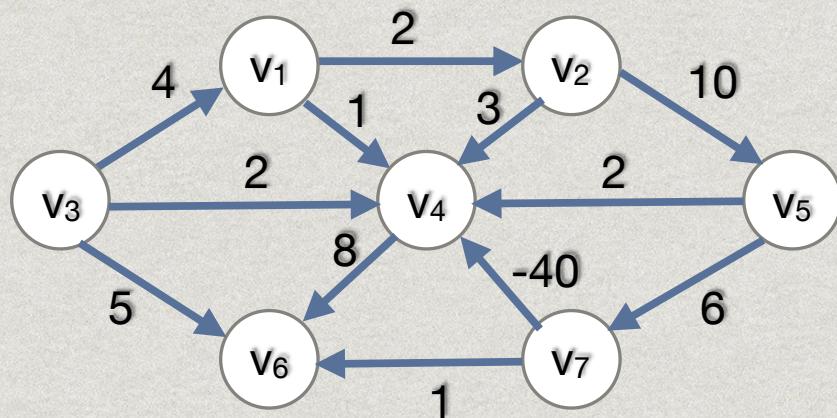
- How Internet Routing works: <https://web.stanford.edu/class/msande91si/www-spr04/readings/week1/InternetWhitepaper.htm>
- <http://www.explainthatstuff.com/internet.html>



Extra Slides



Dijkstra and negative edge costs.

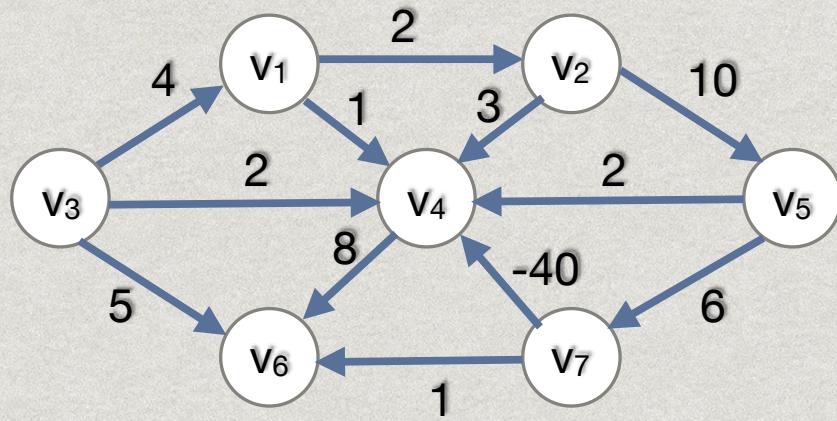


Dijkstra fails with negative edge costs. Once a vertex is declared known (say, v₄), it is possible from some other unknown vertex to create a shorter path to the vertex (say, by eventually looking at v₇ → v₄).

Is there an easy solution?

Dijkstra and negative edge costs.

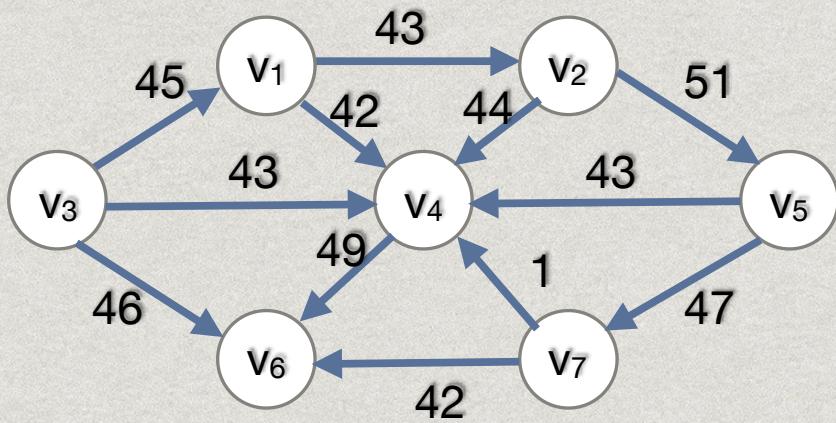
Is there an easy solution?



A naïve approach might be to add a delta (in this case, 41) to all paths, and then apply Dijkstra's algorithm, but this fails because paths with many edges become more weighty than paths with few edges.

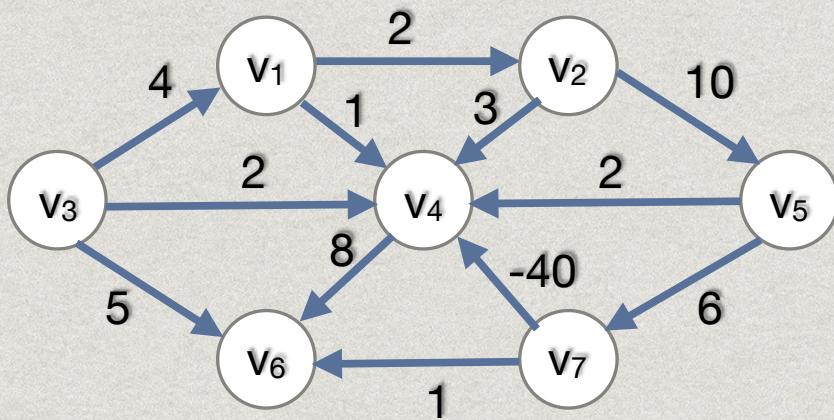
Dijkstra and negative edge costs.

Is there an easy solution?



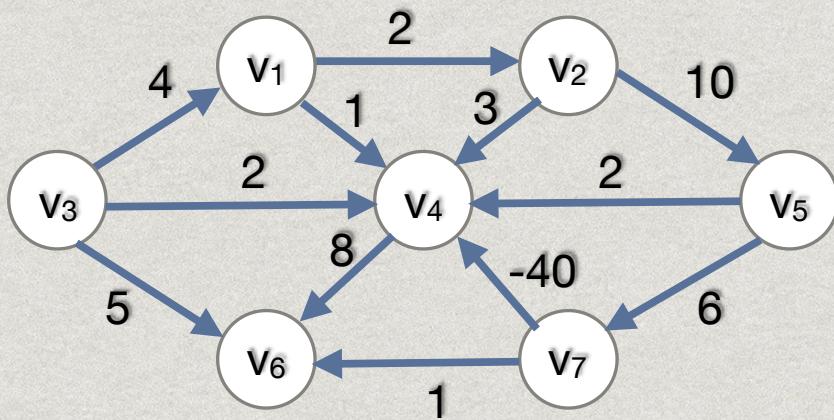
A naïve approach might be to add a delta (in this case, 41) to all paths, and then apply Dijkstra's algorithm, but this fails because paths with many edges become more weighty than paths with few edges.

Dijkstra and negative edge costs.



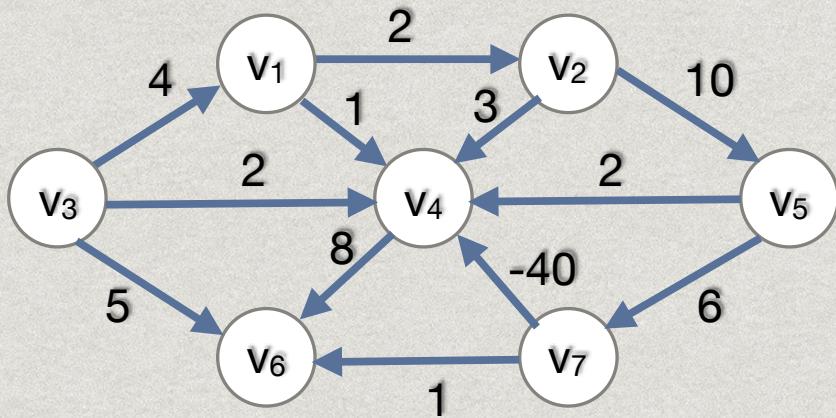
So, there isn't a particularly easy solution. However, we can solve the problem with a combination of the weighted and unweighted algorithms, but at a drastically increased running time cost.

Dijkstra and negative edge costs.



We have to forget about the idea of “known” vertices, since we will have to be able to change our mind if necessary (the greedy algorithm doesn’t work properly).

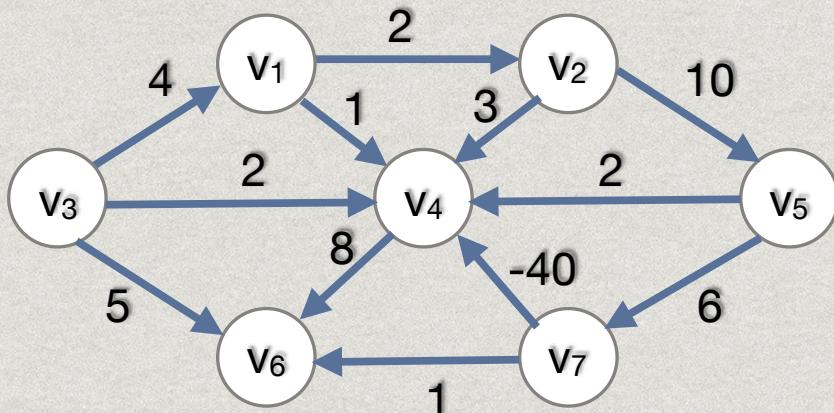
Dijkstra and negative edge costs.



Idea:

1. Place s on a queue
2. At each stage, dequeue a vertex, v. Then, find all vertices, w, adjacent to v, such that:
 $d_w > d_v + \text{cost}_{v,w}$
3. Update d_w and place w on the queue if it isn't already there (set a boolean to indicate presence in the queue)
4. Repeat until the queue is empty.

Dijkstra and negative edge costs.



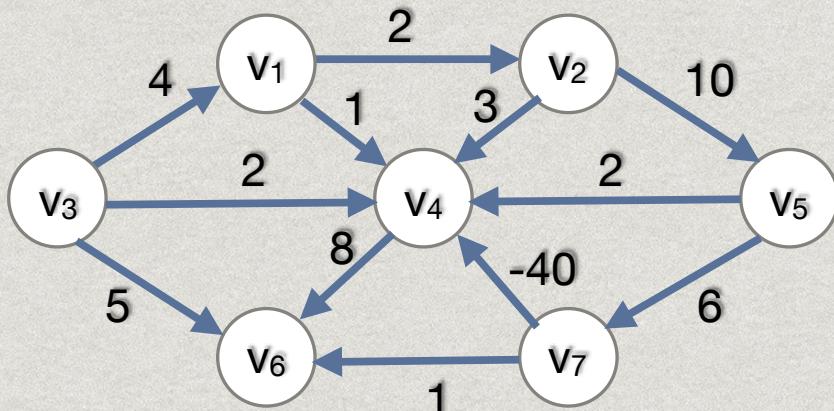
v	in queue?	d _v	p _v
v ₁	Yes	0	0
v ₂	No	INF	0
v ₃	No	INF	0
v ₄	No	INF	0
v ₅	No	INF	0
v ₆	No	INF	0
v ₇	No	INF	0

queue

v ₁					
----------------	--	--	--	--	--

Dequeue v₁ and check weights for v₂ and v₄. Update and place in queue.

Dijkstra and negative edge costs.



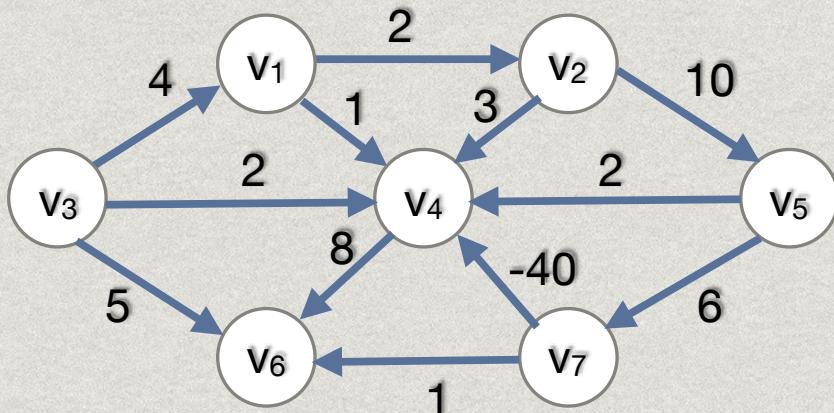
v	in queue?	d_v	p_v
v1	No	0	0
v2	Yes	2	v1
v3	No	INF	0
v4	Yes	1	v1
v5	No	INF	0
v6	No	INF	0
v7	No	INF	0

queue

v2	v4				

Dequeue v₂ and check weights for v₄ and v₅. Update and place in queue.

Dijkstra and negative edge costs.



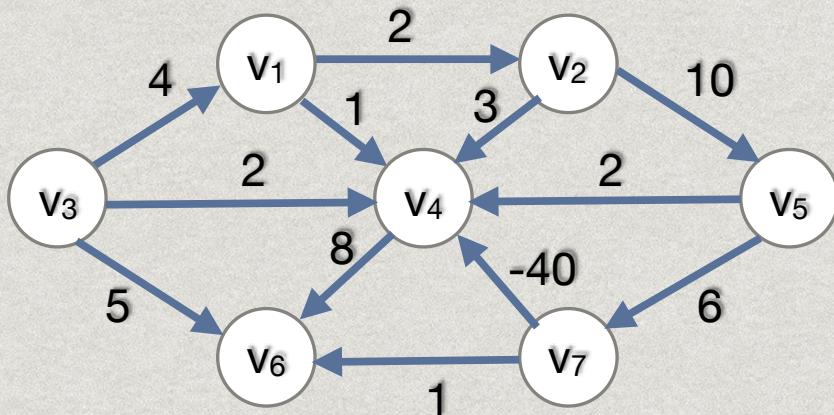
v	in queue?	d_v	p_v
v1	No	0	0
v2	No	2	v1
v3	No	INF	0
v4	Yes	1	v1
v5	Yes	12	v2
v6	No	INF	0
v7	No	INF	0

queue

v4	v5				

Dequeue v4 and check weight for v6. Update and place in queue.

Dijkstra and negative edge costs.



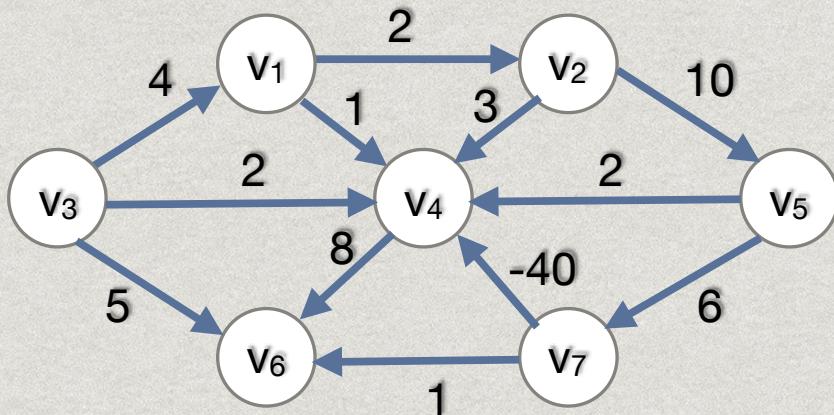
v	in queue?	d_v	p_v
v1	No	0	0
v2	No	2	v1
v3	No	INF	0
v4	No	1	v1
v5	Yes	12	v2
v6	Yes	9	v4
v7	No	INF	0

queue

v5	v6				

Dequeue v₅ and check weights for v₄ and v₇. Update and place in queue.

Dijkstra and negative edge costs.



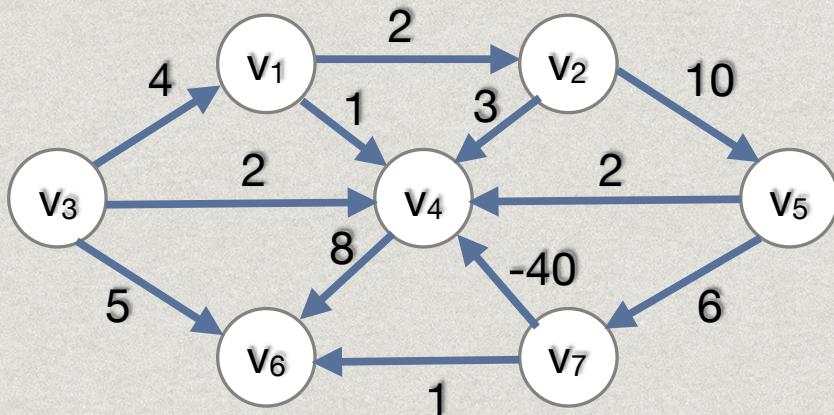
v	in queue?	d _v	p _v
v ₁	No	0	0
v ₂	No	2	v ₁
v ₃	No	INF	0
v ₄	No	1	v ₁
v ₅	No	12	v ₂
v ₆	Yes	9	v ₄
v ₇	Yes	18	v ₅

queue

v ₆	v ₇				

Dequeue v₆ and there aren't any weights to check (v₆ doesn't have any out-going edges).

Dijkstra and negative edge costs.



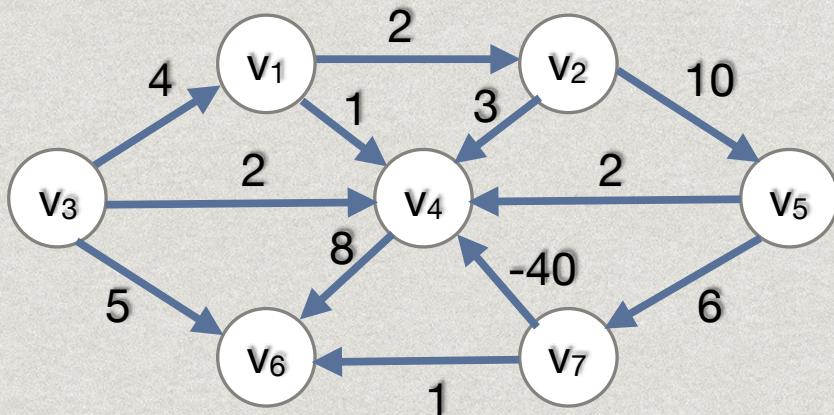
v	in queue?	d _v	p _v
v ₁	No	0	0
v ₂	No	2	v ₁
v ₃	No	INF	0
v ₄	No	1	v ₇
v ₅	No	12	v ₂
v ₆	No	9	v ₄
v ₇	Yes	18	v ₅

queue

v ₇					
----------------	--	--	--	--	--

Dequeue v₇ and check weights for v₄ and v₆. v₄ will get updated and placed back in the queue.

Dijkstra and negative edge costs.



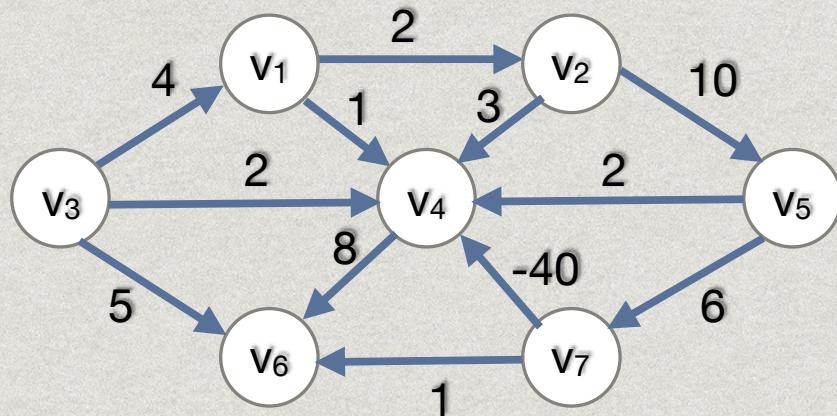
v	in queue?	d_v	p_v
v1	No	0	0
v2	No	2	v1
v3	No	INF	0
v4	Yes	-22	v7
v5	No	12	v2
v6	No	9	v4
v7	No	18	v5

queue

v4					
----	--	--	--	--	--

Dequeue v4 and check weight for v6. v6 will get updated and placed back in the queue.

Dijkstra and negative edge costs.



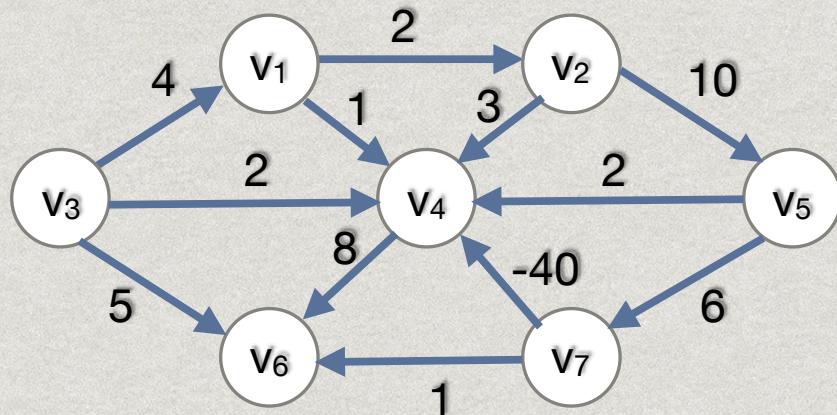
v	in queue?	d _v	p _v
v ₁	No	0	0
v ₂	No	2	v ₁
v ₃	No	INF	0
v ₄	No	-22	v ₇
v ₅	No	12	v ₂
v ₆	Yes	-14	v ₄
v ₇	No	18	v ₅

queue

v ₆					
----------------	--	--	--	--	--

Finally, Dequeue v₆. There aren't any vertices to check (no out-going edges from v₆), and that will empty the queue.

Dijkstra and negative edge costs.



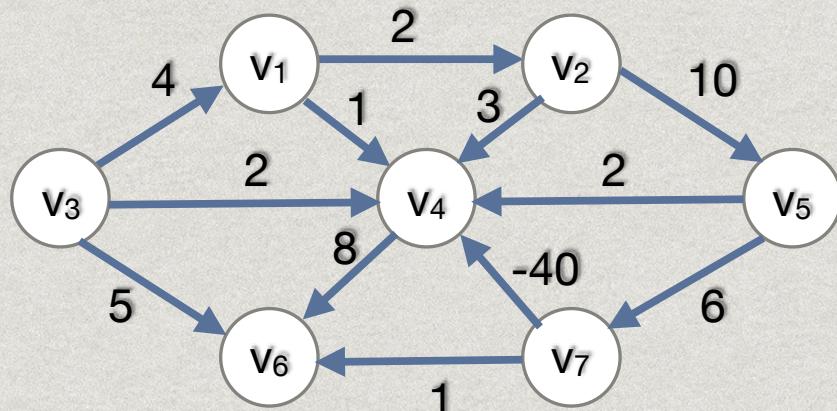
v	in queue?	d_v	p_v
v ₁	No	0	0
v ₂	No	2	v ₁
v ₃	No	INF	0
v ₄	No	-22	v ₇
v ₅	No	12	v ₂
v ₆	No	-14	v ₄
v ₇	No	18	v ₅

queue

--	--	--	--	--	--

Finally, Dequeue v₆. There aren't any vertices to check (no out-going edges from v₆), and that will empty the queue.

Dijkstra and negative edge costs.

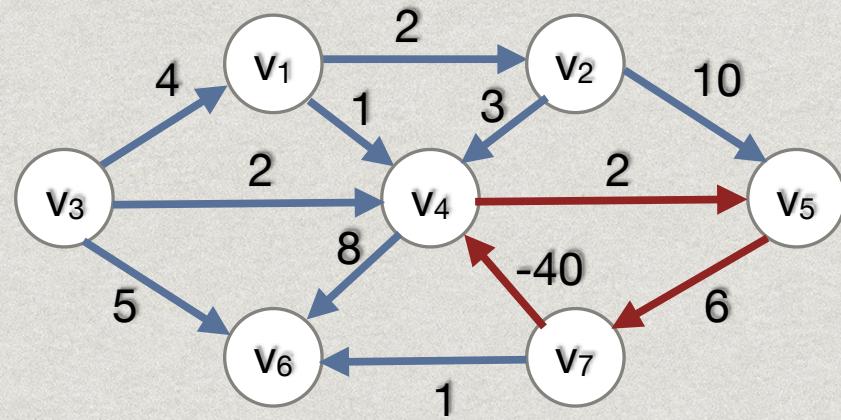


v	in queue?	d _v	p _v
v ₁	No	0	0
v ₂	No	2	v ₁
v ₃	No	INF	0
v ₄	No	-22	v ₇
v ₅	No	12	v ₂
v ₆	No	-14	v ₄
v ₇	No	18	v ₅

Running time?

Each vertex can be dequeued at most $|V|$ times per edge, meaning that the running time is now $O(|E| * |V|)$, which is significantly worse than for the algorithm without negative costs.

Dijkstra and negative edge costs.



What about negative cost cycles?

This will run our queue indefinitely, so we need to make a decision about when to stop.

If you stop after every vertex has been dequeued $|V|+1$ times, you will guarantee termination.